

# COSC 3337 : Data Science I



N. Rizk

College of Natural and Applied Sciences  
Department of Computer Science  
University of Houston

# Hadoop



## A distributed framework for Big Data

Designed to answer the question:

**“How to process big data with reasonable cost and time?”**

# Problem



- How do you scale up applications?
  - Run jobs processing 100's of terabytes of data
  - Takes 11 days to read on 1 computer
- Need lots of cheap computers
  - Fixes speed problem (15 minutes on 1000 computers), but...
  - Reliability problems
    - In large clusters, computers fail every day
    - Cluster size is not fixed
- Need common infrastructure
  - Must be efficient and reliable

# Solution




- Open Source Apache Project
- Hadoop Core includes:
  - Distributed File System - distributes data
  - Map/Reduce - distributes application
- Written in Java
- Runs on
  - Linux, Mac OS/X, Windows, and Solaris
  - Commodity hardware

**The Apache Software Foundation** [/əˈpætʃi/](#) (**ASF**) is an American [nonprofit corporation](#) (classified as a [501\(c\)\(3\) organization](#) in the United States) to support Apache software projects, including the [Apache HTTP Server](#). The ASF was formed from the Apache Group and incorporated on March 25, 1999.

# What is Hadoop?



- **Hadoop:**
  - an open-source software framework that supports data-intensive distributed applications, licensed under the Apache v2 license.
- **Goals / Requirements:**
  - Abstract and facilitate the storage and processing of large and/or rapidly growing data sets
    - Structured and non-structured data
    - Simple programming models
  - High scalability and availability
  - Use commodity (cheap!) hardware with little redundancy
  - Fault-tolerance
  - Move computation rather than data



# What is Hadoop?


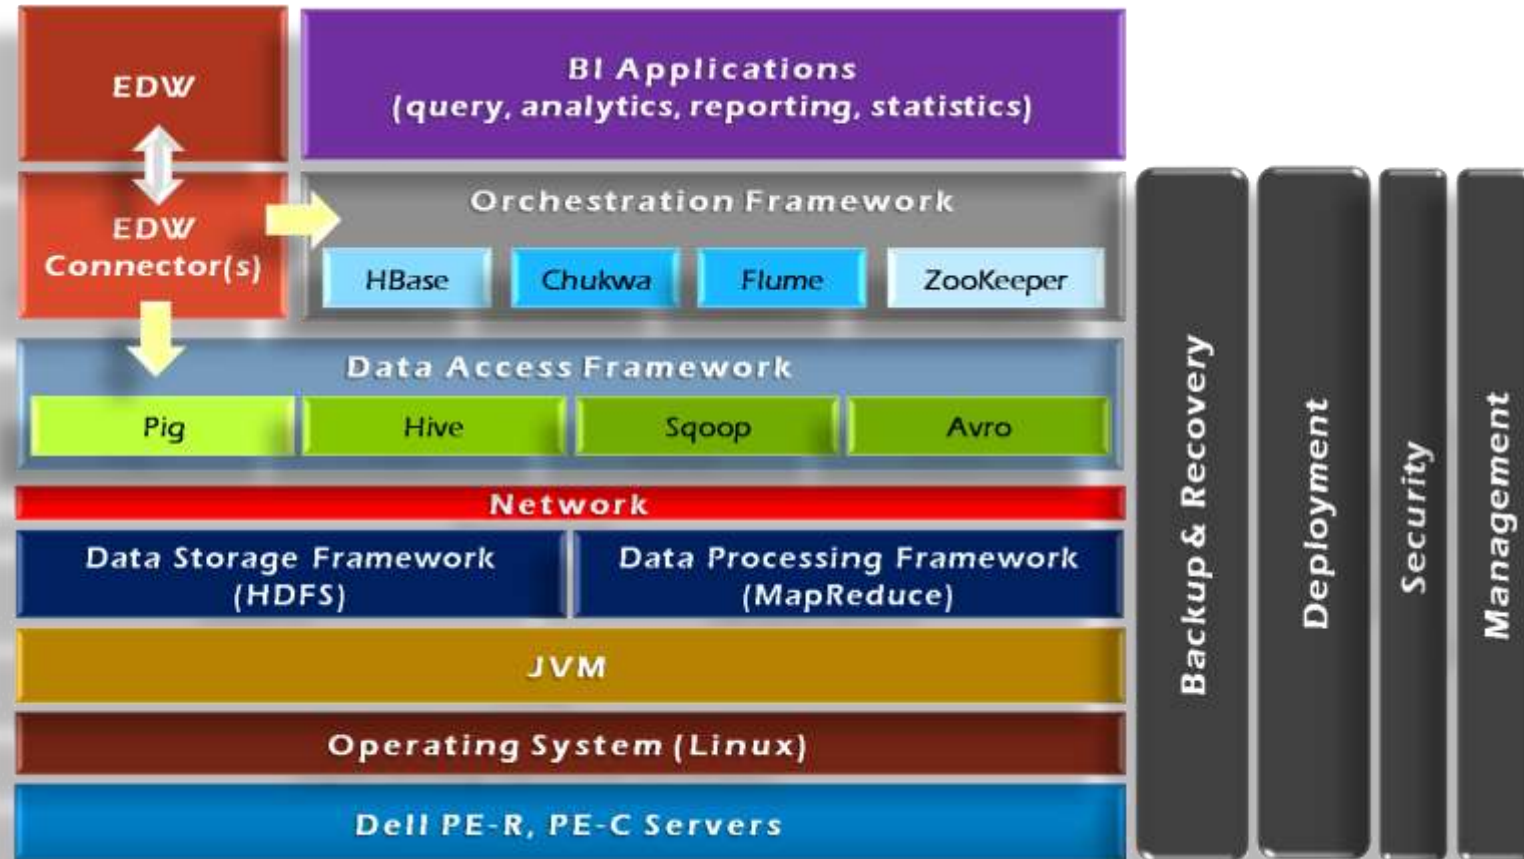


- Apache top level project, **open-source implementation of frameworks** for reliable, scalable, distributed computing and data storage.
- It is a flexible and highly-available **architecture** for large scale computation and data processing on a network of commodity hardware.

# Hadoop Framework Tools



- **2008 - Hadoop Wins Terabyte Sort Benchmark** (sorted 1 terabyte of data in 209 seconds, compared to previous record of 297 seconds)
- 2009 - **Avro and Chukwa** became new members of Hadoop Framework family
- 2010 - Hadoop's **Hbase, Hive and Pig** subprojects completed, adding more computational power to Hadoop framework
- **2011 - ZooKeeper Completed**
- **2013 - Hadoop 1.1.2 and Hadoop 2.0.3 alpha.**
  - Ambari, Cassandra, Mahout have been added



# Hadoop's Developers



Doug Cutting

**2005:** Doug Cutting and Michael J. Cafarella developed Hadoop to support distribution for the [Nutch](#) search engine project.

The project was funded by Yahoo.

**2006:** Yahoo gave the project to Apache Software Foundation.





# Google Origins



2003

## The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung  
Google\*



2004

## MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat  
jeff@google.com, sanjay@google.com

Google, Inc.



2006


## Bigtable: A Distributed Storage System for Structured Data

Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach,  
Mike Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber  
(fay,jeff,sanjay,wilson,deborah,mike,tushar,andy,robert)@google.com  
Google, Inc.


### Abstract

Bigtable is a distributed storage system for managing structured data that is designed to scale to a very large number of nodes across thousands of commodity servers. Many projects at Google store data in Bigtable, including web indexing, Google Earth, and Google Maps. These applications place very different demands on Bigtable, both in terms of data size (from URLs to

achieved scalability and high performance, but Bigtable provides a different interface than such systems. Bigtable does not support a full relational data model; instead, it provides clients with a simple data model that supports dynamic control over data layout and format, allows clients to reason about the locality properties of data represented in the underlying storage. Data is indexed using row and column names that can be arbitrary strings. Bigtable also treats data as uninterpreted bit



# Commodity Hardware Cluster

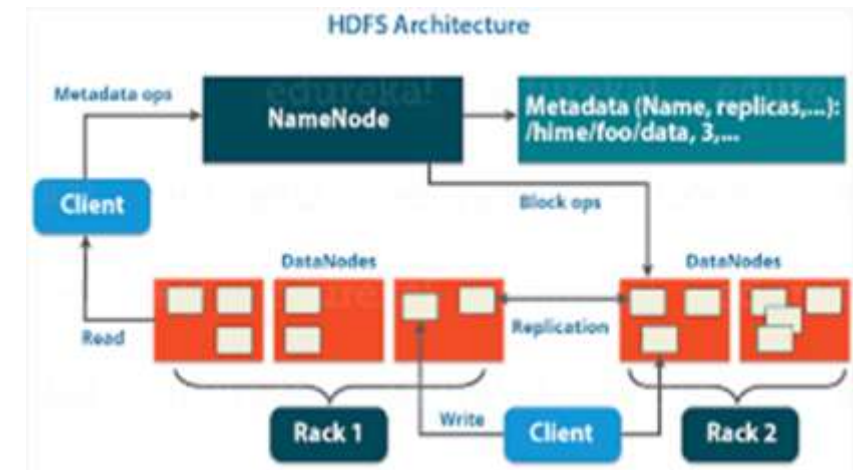



- Typically in 2 level architecture
  - Nodes are commodity PCs
  - 40 nodes/rack
  - Uplink from rack is 8 gigabit
  - Rack-internal is 1 gigabit

# Distributed File System



- Single namespace for entire cluster
  - Managed by a single *namenode*.
  - Files are single-writer and append-only.
  - Optimized for streaming reads of large files.
- Files are broken into large blocks.
  - Typically 128 MB
  - Replicated to several *datanodes*, for reliability
- Client talks to both *namenode* and *datanode*.
  - Data is not sent through the namenode.
  - Throughput of file system scales nearly linearly with the number of nodes.
- Access from Java, C, or command line.

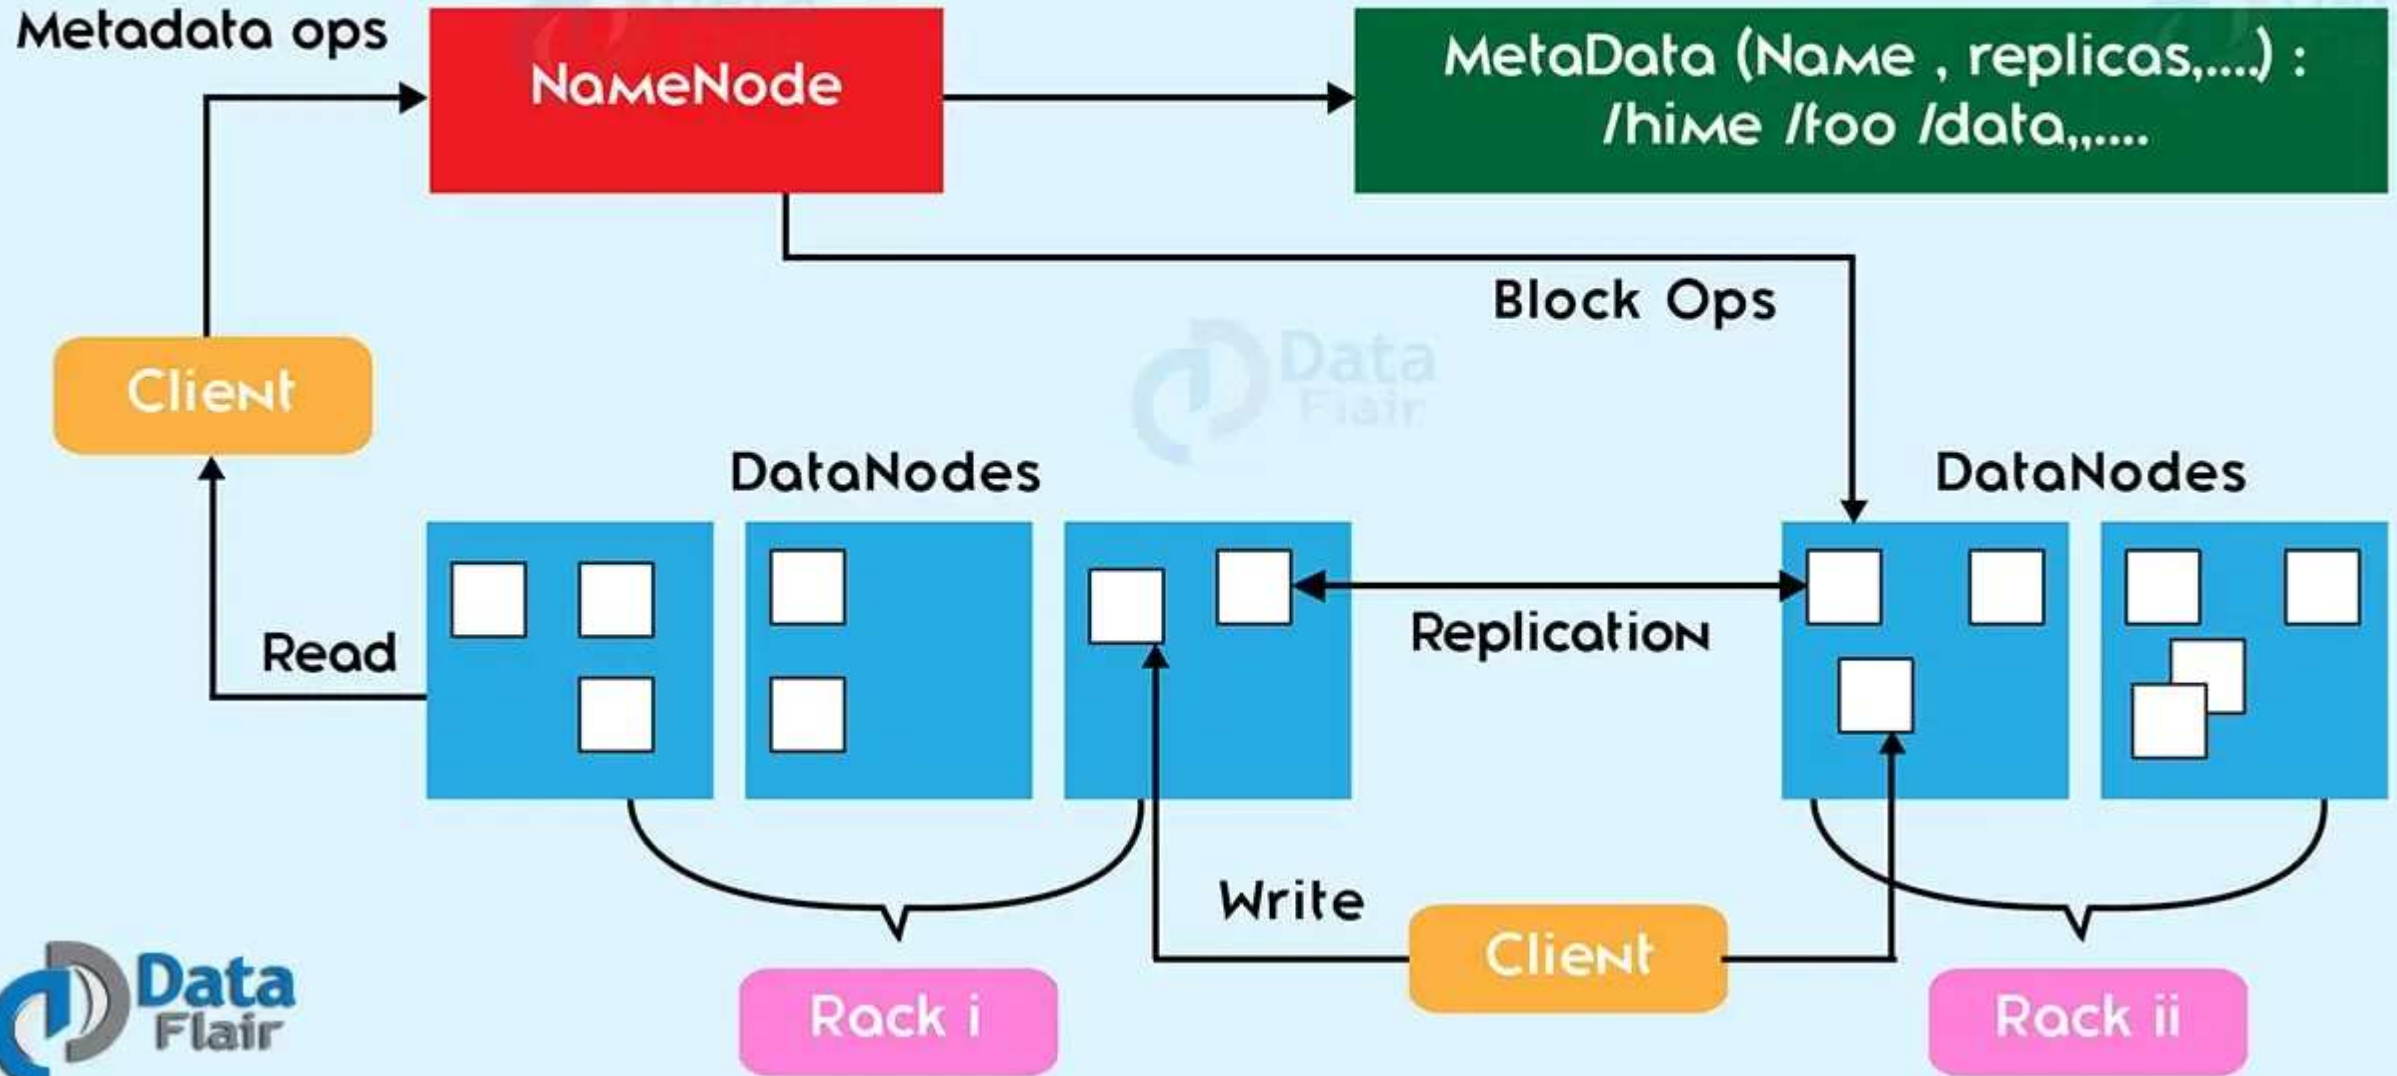


# Block Placement



- Default is 3 replicas, but settable
- Blocks are placed (writes are pipelined):
  - On same node
  - On different rack
  - On the other rack
- Clients read from closest replica
- If the replication for a block drops below target, it is automatically re-replicated.

# HDFS Architecture



# Data Correctness



- Data is checked with CRC32
- File Creation
  - Client computes checksum per 512 byte
  - DataNode stores the checksum
- File access
  - Client retrieves the data and checksum from DataNode
  - If Validation fails, Client tries other replicas
- Periodic Validation

# Map/Reduce

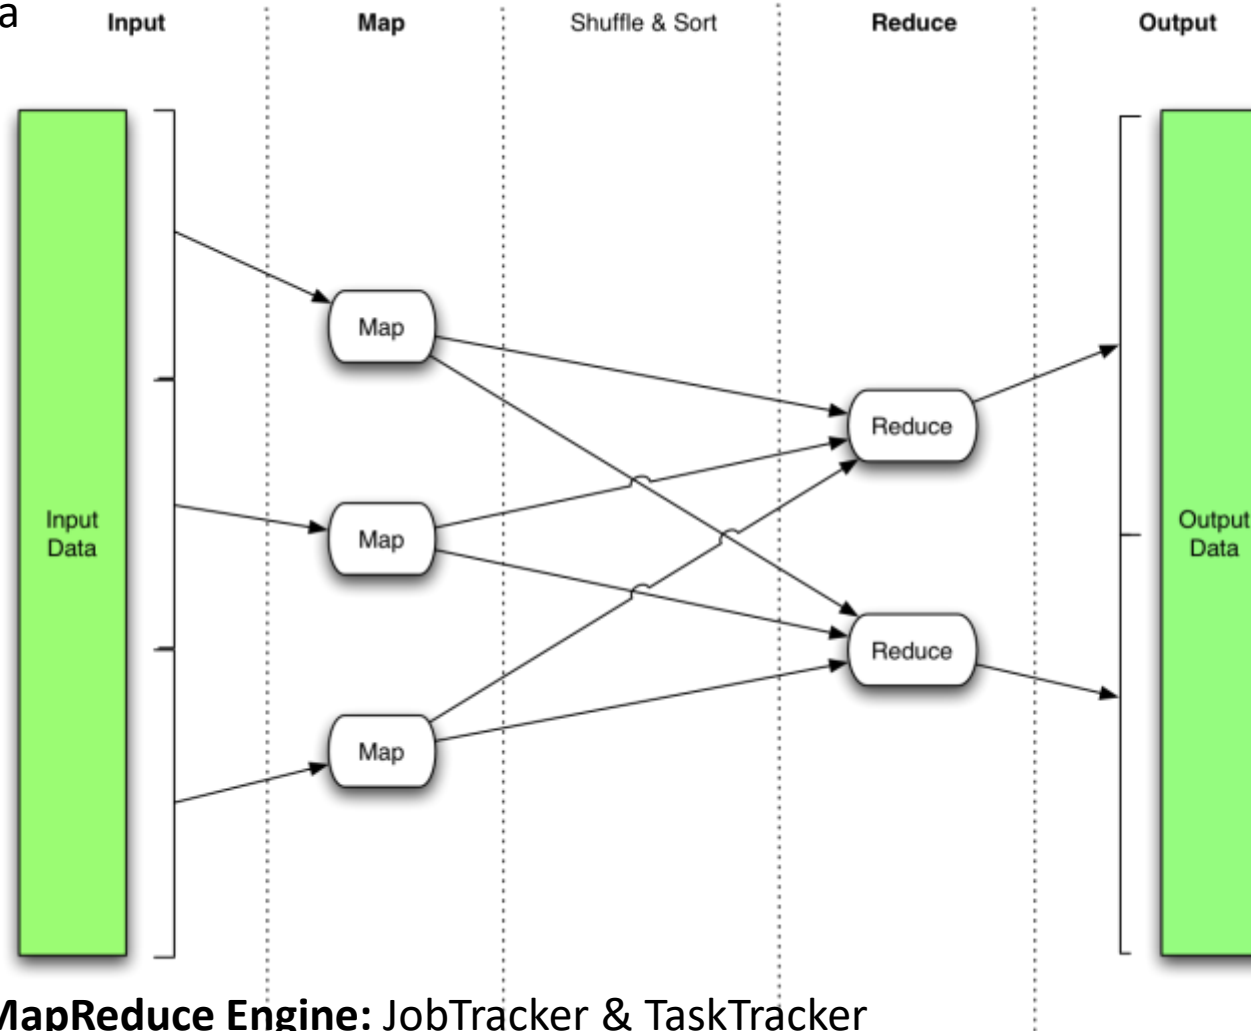


- Map/Reduce is a programming model for efficient distributed computing
- It works like a Unix pipeline:
  - `cat input | grep | sort | uniq -c | cat > output`
  - **Input** | **Map** | Shuffle & Sort | **Reduce** | **Output**
- Efficiency from
  - Streaming through data, reducing seeks
  - Pipelining
- A good fit for a lot of applications
  - Log processing
  - Web index building

# Map/Reduce Dataflow



- JobTracker splits up data into smaller tasks(“Map”) and sends it to the TaskTracker process in each node


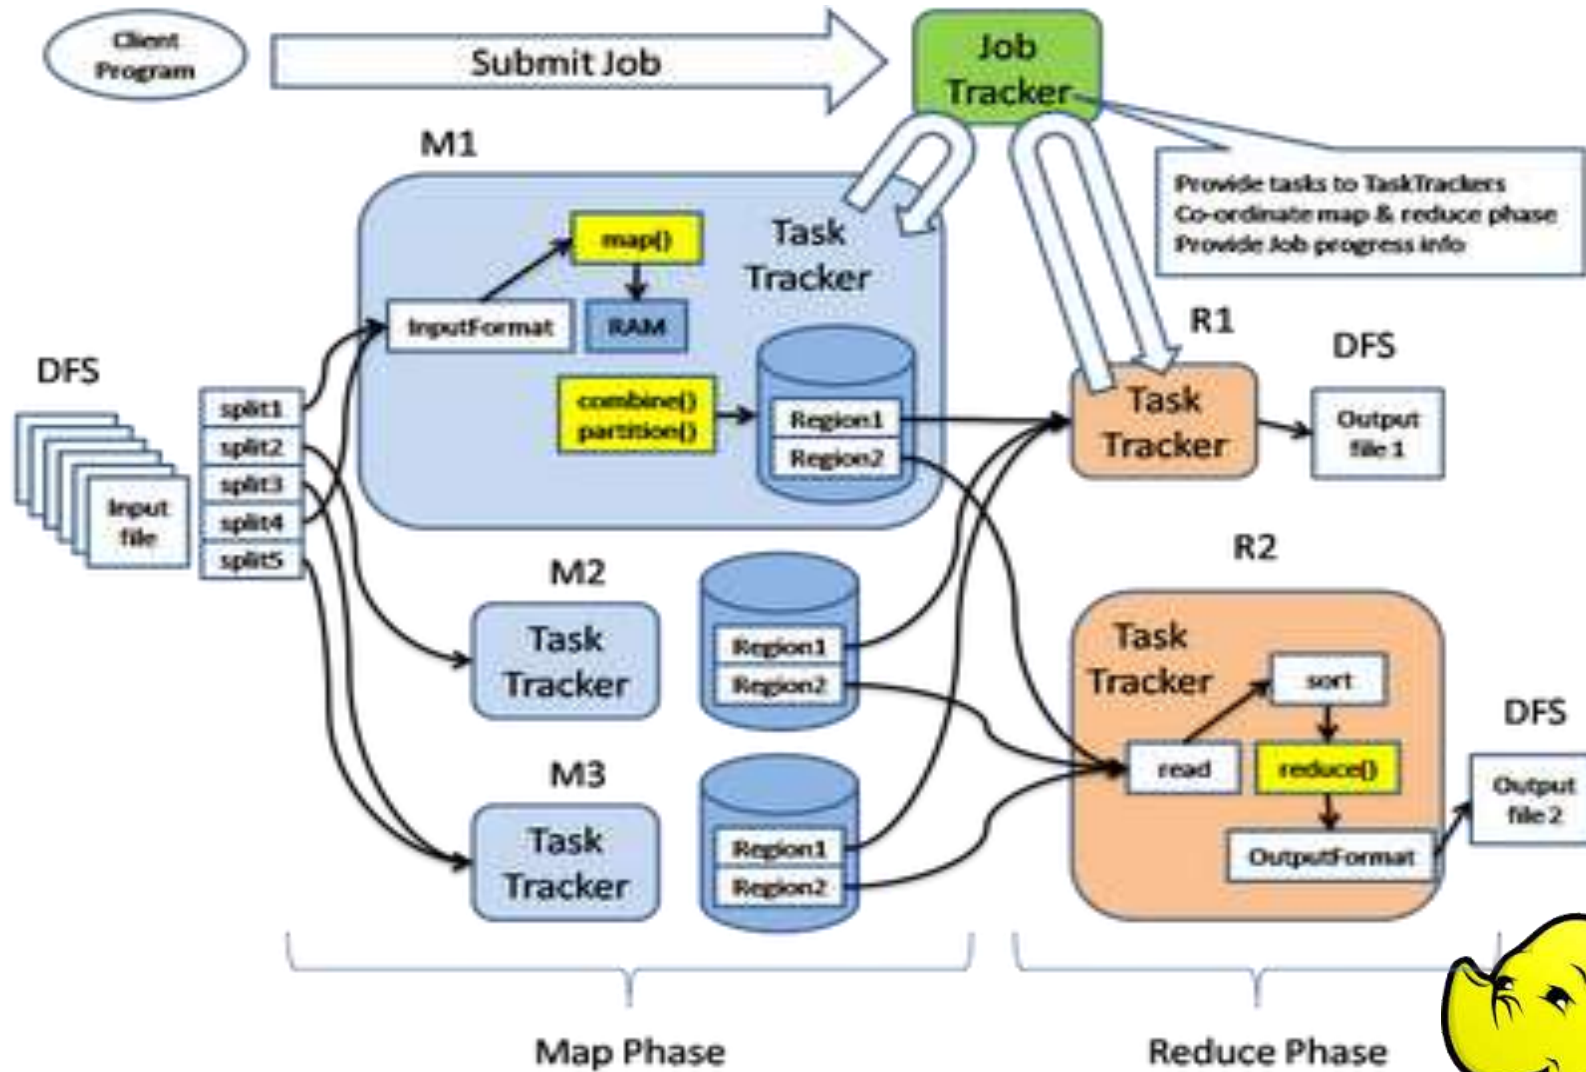


- TaskTracker reports back to the JobTracker node and reports on job progress, sends data (“Reduce”) or requests new jobs

**MapReduce Engine:** JobTracker & TaskTracker



# Hadoop's Architecture: MapReduce Engine



# Map/Reduce features



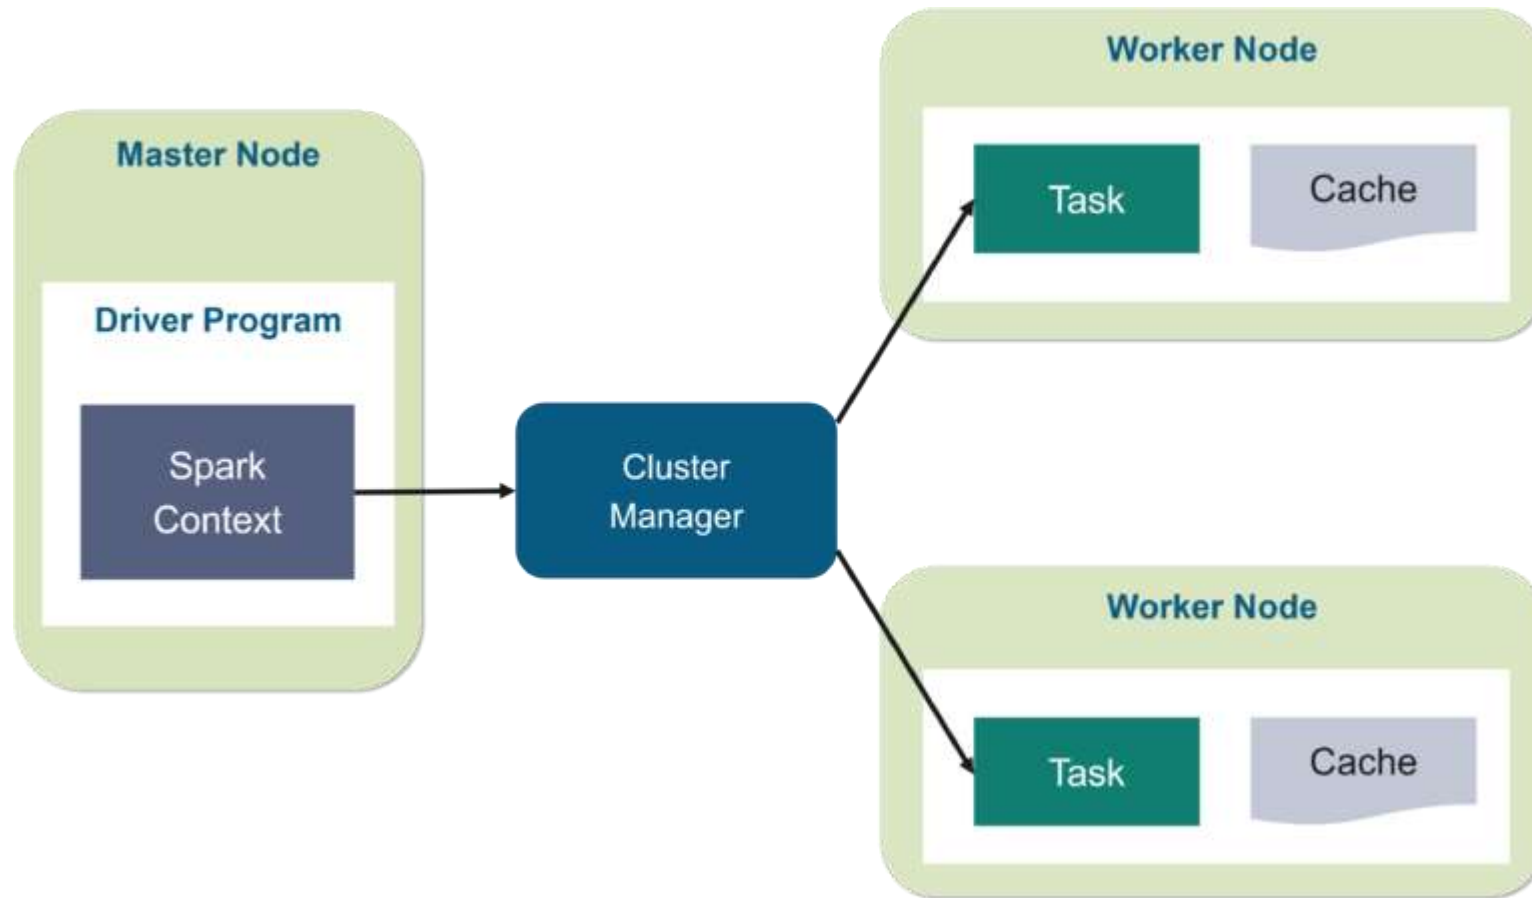
- Java and C++ APIs
  - In Java use Objects, while in C++ bytes
- Each task can process data sets larger than RAM
- Automatic re-execution on failure
  - In a large cluster, some nodes are always slow or flaky
  - Framework re-executes failed tasks
- Locality optimizations
  - Map-Reduce queries HDFS for locations of input data
  - Map tasks are scheduled close to the inputs when possible

# Hadoop clusters



- ~20,000 machines running Hadoop
- clusters are 2000 nodes
- Several petabytes of user data (compressed, unreplicated)
- Run hundreds of thousands of jobs every month

# Spark Architecture



# SPARK



Spark is a cluster – computing framework. It competes with MapReduce and the entire Hadoop ecosystem. Spark doesn't have its own distributed filesystem, but it uses the HDFS.

## Polyglot

Spark provides high-level APIs in Java, Scala, Python and R. Spark code can be written in any of these languages. It also provides a shell in Scala and Python.

## Speed

Spark runs up to 100 times faster than Hadoop MapReduce for large-scale data processing. It is also able to achieve this speed through controlled partitioning.

## PowerfulCaching

Simple programming layer provides powerful caching and disk persistence capabilities.

## Deployment

It can be deployed through *Mesos, Hadoop via YARN, or Spark's own cluster manager.*

## Real-Time

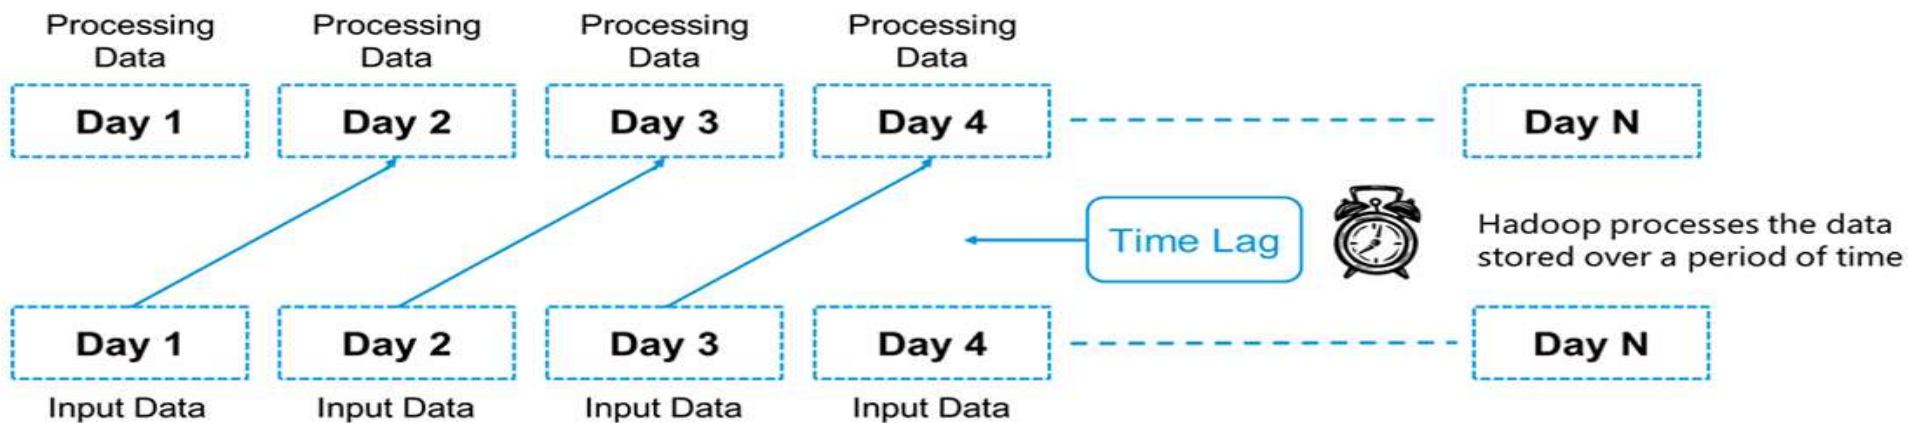
It offers Real-time computation & low latency because of *in-memory computation.*




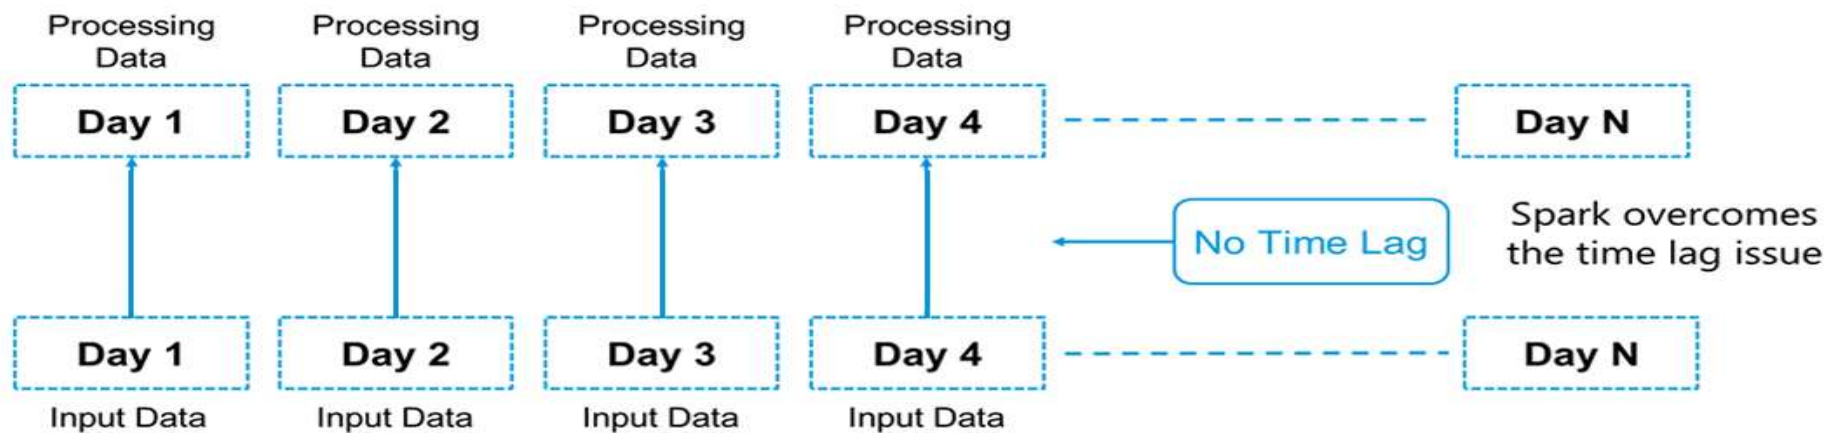




## Processing Data Using MapReduce




## Real Time Processing in Spark



# Three main applications of Hadoop:




- Advertisement (Mining user behavior to generate recommendations)
- Searches (group related documents)
- Security (search for uncommon patterns)
- Hadoop is in use at most organizations that handle big data:
  - Yahoo!
  - Facebook
  - Amazon
  - Netflix
  - Etc...



# Hadoop in the Wild



- Non-realtime large dataset computing:
  - NY Times was dynamically generating PDFs of articles from 1851-1922
  - Wanted to pre-generate & statically serve articles to improve performance
  - Using Hadoop + MapReduce running on EC2 / S3, converted 4TB of TIFFs into 11 million PDF articles in 24 hrs






# Hadoop in the Wild: Facebook Messages




- Design requirements:
  - Integrate display of email, SMS and chat messages between pairs and groups of users
  - Strong control over who users receive messages from
  - Suited for production use between **500 million people immediately after launch**
  - Stringent latency & uptime requirements



# Hadoop in the Wild




- Facebook's solution
  - Hadoop + HBase as foundations
  - Improve & adapt HDFS and HBase to scale to FB's workload and operational considerations
    - Major concern was availability: NameNode is SPOF & failover times are at least 20 minutes
    - Proprietary "AvatarNode": eliminates SPOF, makes HDFS safe to deploy even with 24/7 uptime requirement
    - Performance improvements for realtime workload: RPC timeout. Rather fail fast and try a different DataNode



# Hadoop in the Wild



- Classic alternatives
  - These requirements typically met using large MySQL cluster & caching tiers using Memcached
  - Content on HDFS could be loaded into MySQL or Memcached if needed by web tier
- Problems with previous solutions
  - MySQL has low random write throughput... BIG problem for messaging!
  - Difficult to scale MySQL clusters rapidly while maintaining performance
  - MySQL clusters have high management overhead, require more expensive hardware



# Hadoop Highlights



- Distributed File System
- Fault Tolerance
- Open Data Format
- Flexible Schema
- Queryable Database

# Why use Hadoop?

- Need to process Multi Petabyte Datasets
- Data may not have strict schema
- Expensive to build reliability in each application
- Nodes fails everyday
- Need common infrastructure
- Very Large Distributed File System
- Assumes Commodity Hardware
- Optimized for Batch Processing
- Runs on heterogeneous OS