

# COSC 3337 : Data Science I



N. Rizk

College of Natural and Applied Sciences  
Department of Computer Science  
University of Houston

SUCCESS

HARD WORK  
PERSISTENCE  
LATE NIGHTS  
REJECTIONS  
SACRIFICES  
DISCIPLINE  
CRITICISM  
DOUBTS  
FAILURE  
RISKS

# Outline



[Getting started](#)

[Explore dataset content](#)

[Inspect visually](#)

[Run clustering](#)

[Assess clustering quality](#)

# Getting started



Setting up python environment

At first, we need to set up our environment.

```
pip install numpy pandas matplotlib sklearn
```

```
1 # We need pandas to import and use .csv files
2 import pandas as pd

4 # numpy is used for various numeric operations
5 import numpy as np
6
7 # matplotlib lets us draw nice plots
8 import matplotlib.pyplot as plt
9
10 # We will use several modules of sklearn package
11 # that is a "swiss knife" ML toolset for python
12 from sklearn import cluster, metrics, decomposition
```

# Get the Iris dataset

Example data for clustering



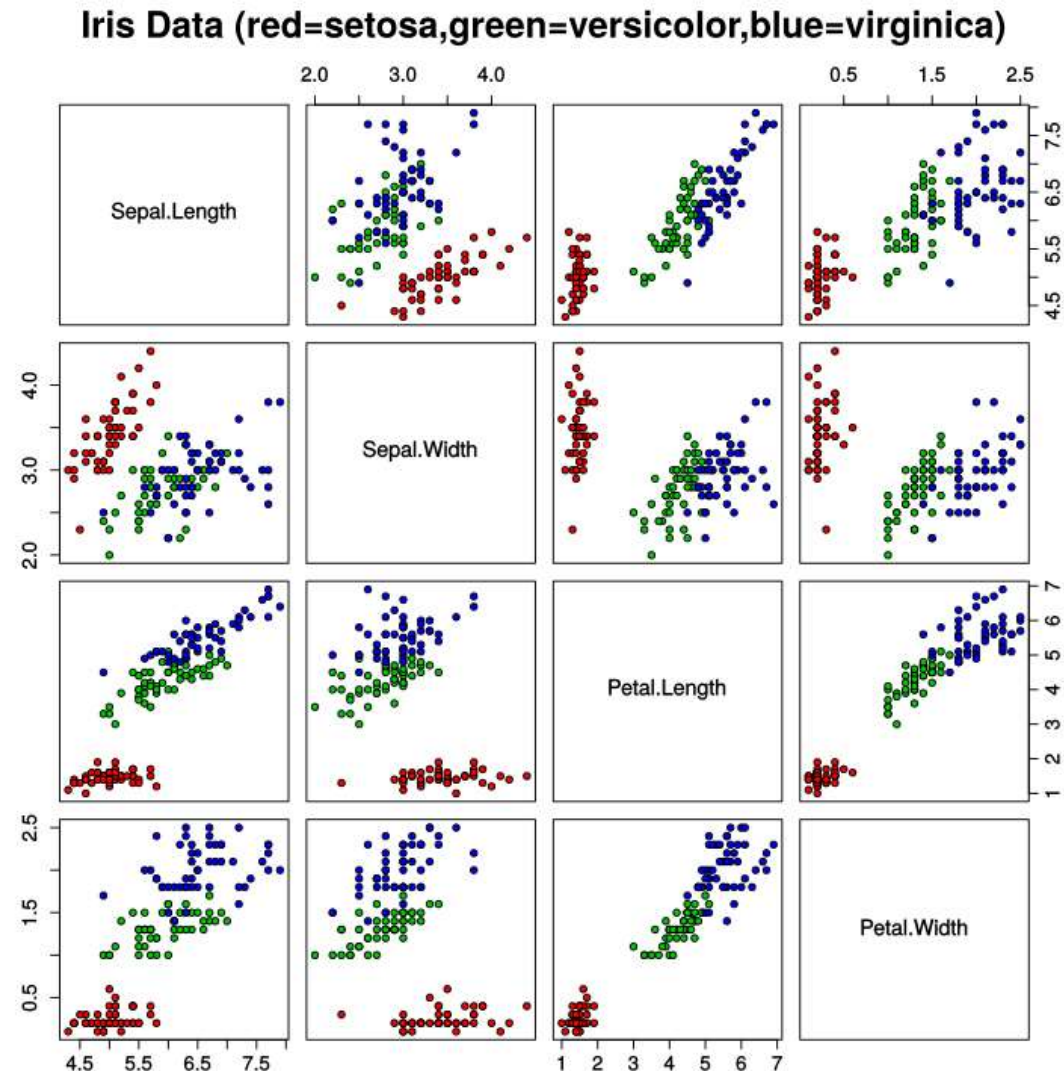
The Iris flower data set is a multivariate data set introduced by the British statistician and biologist Ronald Fisher in 1936.

You can read more about it on Wikipedia's page

[https://en.wikipedia.org/wiki/Iris\\_flower\\_data\\_set](https://en.wikipedia.org/wiki/Iris_flower_data_set)

You can find the .csv file easily; I got it from

<https://raw.githubusercontent.com/vincentarelbundock/Rdatasets/master/csv/datasets/iris.csv>



# Get the Iris dataset



Now, we are ready to load the .csv file into the interpreter using pandas package.

```
1 # Using pandas library to load csv file
2 # into a DataFrame object
3 ourData = pd.read_csv("iris.csv")
4
5 # Created object `ourData` contains
6 # a funny and famous data set of flowers.
7 # At first, we need to explore,
8 # what is in this data set.
9 # For this, pandas package provides
10 # several very useful functions
```



# Explore the Iris dataset

Trying useful functions from `pandas`

The first useful command is `head()`

It returns first several lines of a dataframe. Very useful to get an idea of what kind of data you have!

```
ourData . head()
```

#	Unnamed: 0	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
# 0	1	5.1	3.5	1.4	0.2	setosa
# 1	2	4.9	3.0	1.4	0.2	setosa
# 2	3	4.7	3.2	1.3	0.2	setosa
# 3	4	4.6	3.1	1.5	0.2	setosa
# 4	5	5.0	3.6	1.4	0.2	setosa

# Explore the Iris dataset

Trying useful functions from `pandas`

Second command, `info()`, can give a more precise information on what data types we operate on.

- ▶ The dataset has 150 entries
- ▶ Sepal and petal parameters are floating-point numbers
- ▶ Species is recognized as object type – in fact, this is a text

```
# <class 'pandas.core.frame.DataFrame'>
# Int64Index: 150 entries, 0 to 149
# Data columns (total 6 columns):
# Unnamed: 0      150 non-null int64
# Sepal.Length    150 non-null float64
# Sepal.Width     150 non-null float64
# Petal.Length    150 non-null float64
# Petal.Width     150 non-null float64
# Species         150 non-null object
# dtypes: float64(4), int64(1), object(1)
# memory usage: 8.2+ KB
```



# Explore the Iris dataset

Trying useful functions from `pandas`

The last command is more advanced: `ourData.describe()`.

`ourData.describe()`

This command gives some statistical information about numeric values in your dataset.

It is useful to understand what is the range of your values.

```

#      Unnamed: 0  Sepal.Length Sepal.Width  Petal.Length Petal.Width
# count  150.00000  150.000000  150.000000  150.000000  150.000000
# mean    75.500000    5.843333    3.057333    3.758000    1.199333
# std     43.445368    0.828066    0.435866    1.765298    0.762238
# min      1.000000    4.300000    2.000000    1.000000    0.100000
# 25%     38.250000    5.100000    2.800000    1.600000    0.300000
# 50%     75.500000    5.800000    3.000000    4.350000    1.300000
# 75%    112.750000    6.400000    3.300000    5.100000    1.800000
# max    150.000000    7.900000    4.400000    6.900000    2.500000

```



# Explore the Iris dataset

Trying useful functions from `pandas`

But what about the last column `Species`?

We've seen these are textual values. Now, let us see what kinds of values does this column have?

```
1 # Get unique values in the column
2 species = ourData.Species.unique()
3
4 print(species)

# ['setosa' 'versicolor' 'virginica']
```

# Draw some plots



Using matplotlib

Let's try to plot something!

```
1 # plot different species in different colours
2 colors = ['g', 'r', 'b', 'c', 'm', 'k']
3 species_dict = dict(zip(species, colors))
4 plt.scatter( ourData['Sepal.Length']
5              , ourData['Sepal.Width']
6              , c=ourData['Species']
7              .apply(lambda x: species_dict[x])
8              )
9 # use either
10 # plt.show() # to draw on-screen
11 # or
12 # plt.savefig("filename") # to save picture
13 plt.savefig("real-labels.png")
```

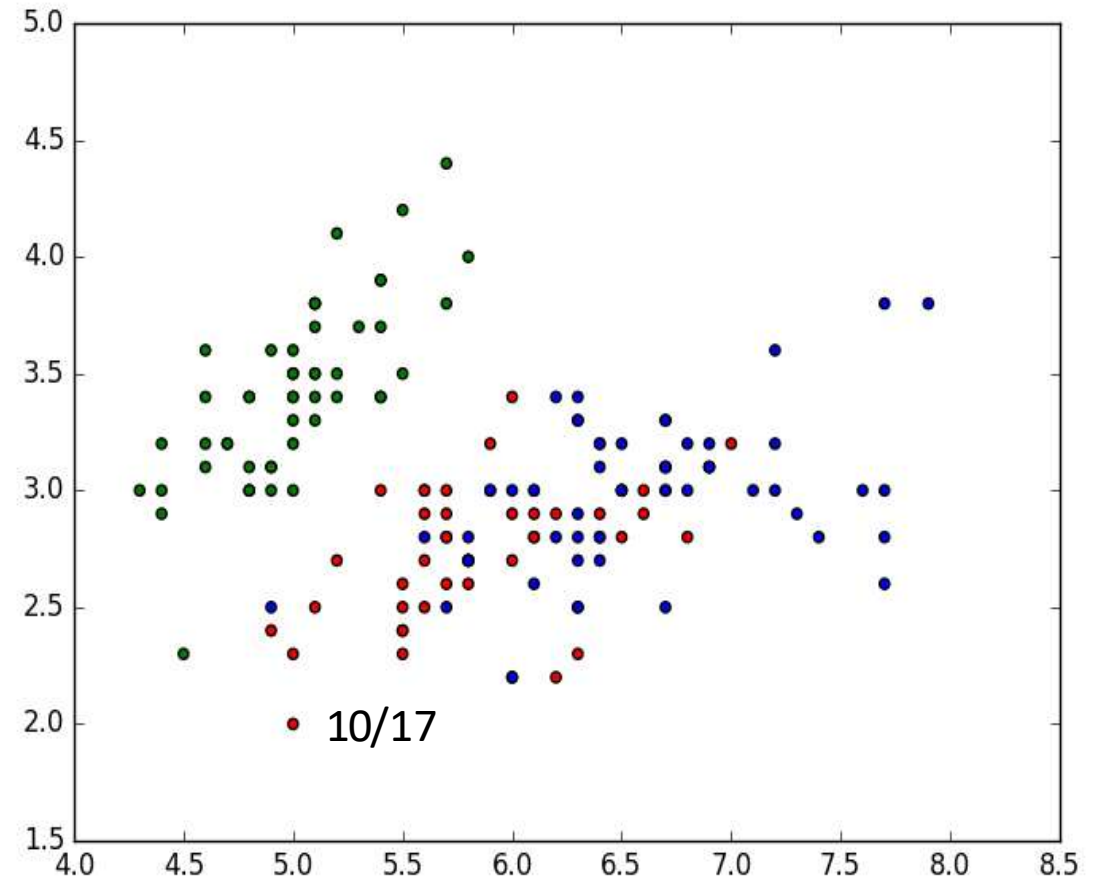
In this example I map names of the species onto colors.  
Play with this code a little bit changing the columns to plot. This gives  
you an insight how the data looks like.

# Draw some plots

Using `matplotlib`



...And here is how the result looks like





# Prepare data

Split input and labels

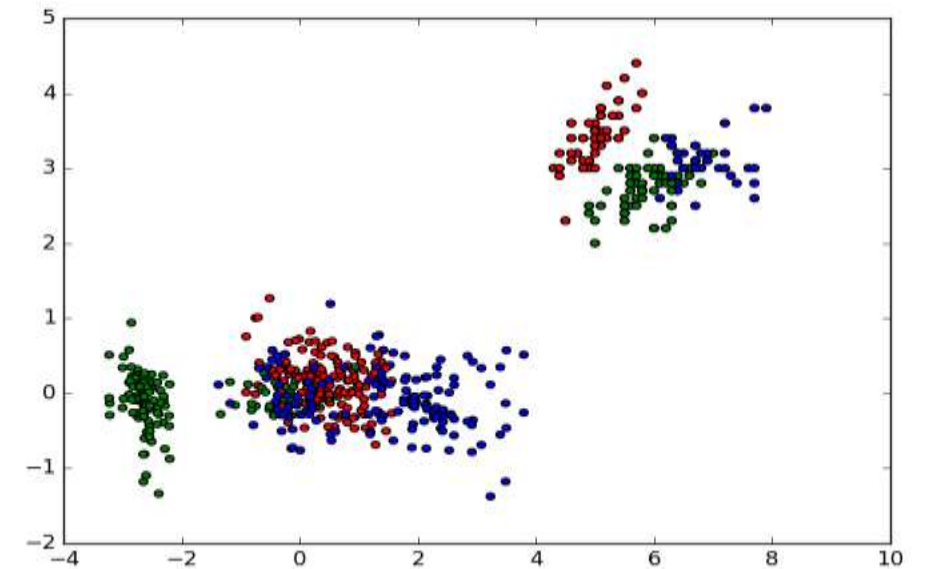
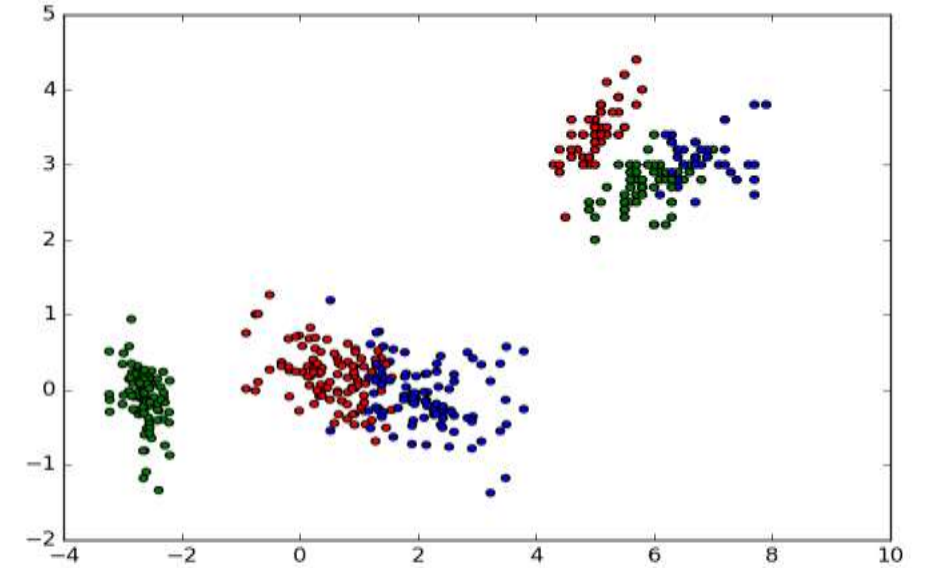
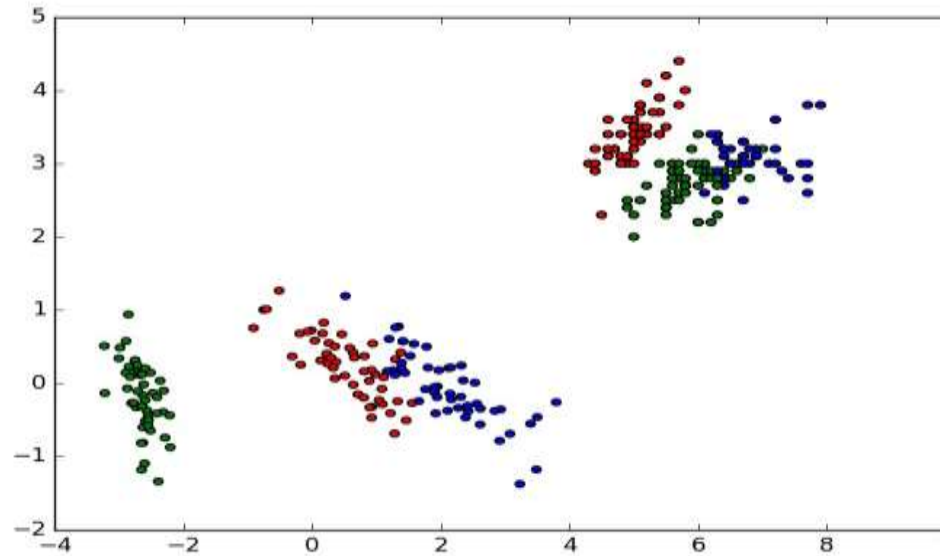
```
1 # Here is our data , but without species information
2 ourDataNoLabels
3 = pd . DataFrame( ourData
4                   , columns=[ 'Sepal . Length '
5                               , 'Sepal . Width '
6                               , 'Petal . Length '
7                               , 'Petal . Width '
8                               ]
9                   )
```

Copy the data to a new variable `ourDataNoLabels`.  
We will play like we don't know the labels and want to estimate them.

# Draw some plots

Using matplotlib

...And here are the plots



# Principal Component analysis

Using `sklearn.decomposition`

This code performs PCA (lines 3-7).  
(Most of the code is just to draw three plots).

Principal component transform is a procedure that rotates point space in such a way that points variance is the biggest along X-axis, the second is along Y-axis, and so on.

The method is simple but very powerful for data exploration. If data have many dimensions (i.e. 50) it is very convenient to look only at 2-5 most significant dimensions. Otherwise visual inspection of the data would be too difficult.

```
1 # Run Principal Component Analysis
2 # to get more understanding how our data looks like
3 ourDataReduced
4     = decomposition
5       .PCA(n_components=3)
6       .fit_transform(ourDataNoLabels)
7 plt.scatter( ourDataReduced[:,0]
8              , ourDataReduced[:,1]
9              , c=ourData['Species']
10                .apply(lambda x: species_dict[x]))
11 plt.savefig("real-pca-1-2.png")
12 plt.scatter( ourDataReduced[:,0]
13              , ourDataReduced[:,2]
14              , c=ourData['Species']
15                .apply(lambda x: species_dict[x]))
16 plt.savefig("real-pca-1-3.png")
17 plt.scatter( ourDataReduced[:,1]
18              , ourDataReduced[:,2]
19              , c=ourData['Species']
20                .apply(lambda x: species_dict[x]))
21 plt.savefig("real-pca-2-3.png")
```

# Principal Component analysis



Using `sklearn.decomposition`

This code performs PCA (lines 3-7).  
(Most of the code is just to draw three plots).

Principal component transform is a procedure that rotates point space in such a way that points variance is the biggest along X-axis, the second is along Y-axis, and so on.

The method is simple but very powerful for data exploration. If data have many dimensions (i.e. 50) it is very convenient to look only at 2-5 most significant dimensions. Otherwise visual inspection of the data would be too difficult.

```
1 # Run Principal Component Analysis
2 # to get more understanding how our data looks like
3 ourDataReduced
4 = decomposition
5 .PCA(n_components=3)
6 .fit_transform(ourDataNoLabels)
7 plt.scatter(ourDataReduced[:,0]
8             , ourDataReduced[:,1]
9             , c=ourData['Species']
10             .apply(lambda x: species_dict[x]))
11 plt.savefig("real-pca-1-2.png")
12 plt.scatter(ourDataReduced[:,0]
13             , ourDataReduced[:,2]
14             , c=ourData['Species']
15             .apply(lambda x: species_dict[x]))
16 plt.savefig("real-pca-1-3.png")
17 plt.scatter(ourDataReduced[:,1]
18             , ourDataReduced[:,2]
19             , c=ourData['Species']
20             .apply(lambda x: species_dict[x]))
21 plt.savefig("real-pca-2-3.png")
```



# K-means clustering

Using `sklearn.cluster`



Finally, we are ready for some clustering!

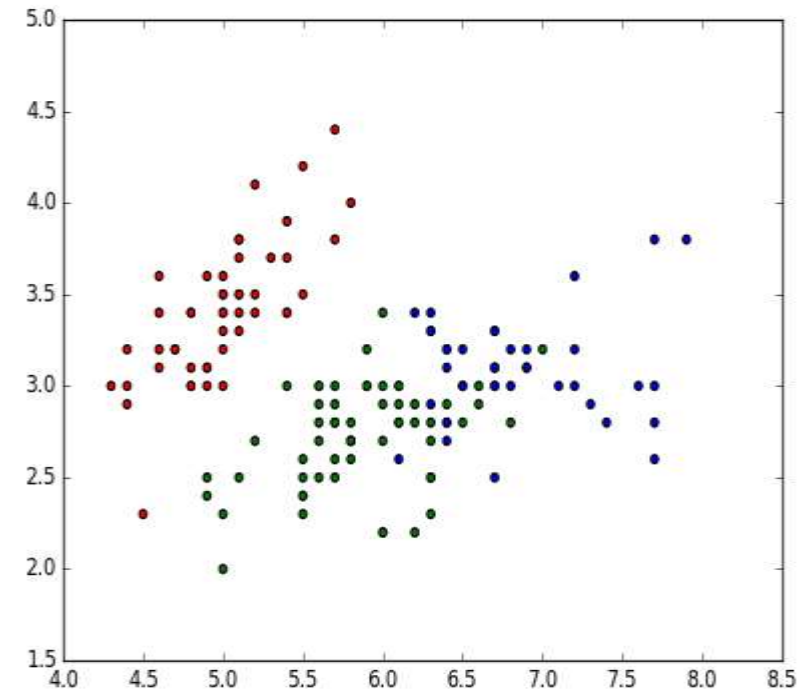
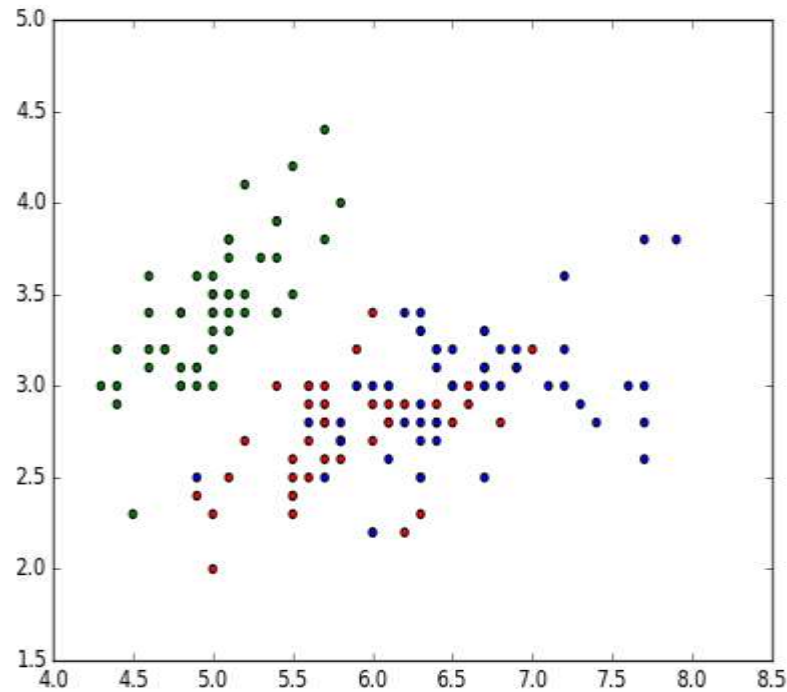
```
1 # We know that there must be 3 different clusters
2 #                               - species of flowers.
3 # So let's give this hint to the algorithm
4 kmeans = cluster.KMeans(3)
5           .fit(np.array(ourDataNoLabels))
6 foundLabels
7   = pd.DataFrame( kmeans.labels_
8                   , columns=['K-means clusters'])
9
10
11 plt.scatter( ourData['Sepal.Length']
12              , ourData['Sepal.Width']
13              , c=foundLabels['K-means clusters']
14              .apply(lambda x: colors[x])
15              )
16 plt.savefig("predicted-labels.png")
```

# K-means clustering

Using `sklearn.cluster`



Let's compare plots...



...would you say this is a good result?

# Assess clustering quality

Using pandas



A very good way to assess performance of an unsupervised clustering algorithm is to look at co-occurrence tables.

```
1 mat = pd.crosstab( ourData['Species']  
2                    , foundLabels['K-means clusters']  
3                    )  
4 print(mat)
```

Package `pandas` provides a special function `crosstab()` that calculates how many times a value from one column occurs together with a value from another column.

```
# K-means clusters    0    1    2  
# Species  
# setosa              0   50    0  
# versicolor         48    0    2  
# virginica          14    0   36
```

So, what would be a conclusion now?

# Bonus: DBSCAN

Using `sklearn.cluster`



This code does everything the same way as KMeans clustering example.

Try it! And compare the results.  
Which algorithm performs better?

```
1 dbscan = cluster.DBSCAN()  
2         .fit(np.array(ourDataNoLabels))  
3 foundLabels['DBSCAN clusters']  
4     = pd.DataFrame( dbscan.labels_  
5                     , columns=['DBSCAN clusters'])  
6 plt.scatter( ourData['Sepal.Length']  
7             , ourData['Sepal.Width']  
8             , c=foundLabels['DBSCAN clusters']  
9             .apply(lambda x: colors[x]))  
10 plt.savefig("predicted-labels2.png")  
11 mat = pd.crosstab( ourData['Species']  
12                  , foundLabels['DBSCAN clusters'])  
13 print(mat)
```