

# COSC 3337 : Data Science I



N. Rizk

College of Natural and Applied Sciences  
Department of Computer Science  
University of Houston

# Focus



Machine Learning

Supervised Learning  
Develop predictive models  
based on both input and  
output data  
(categorical /continuous)

Classification

Output data  
categorical

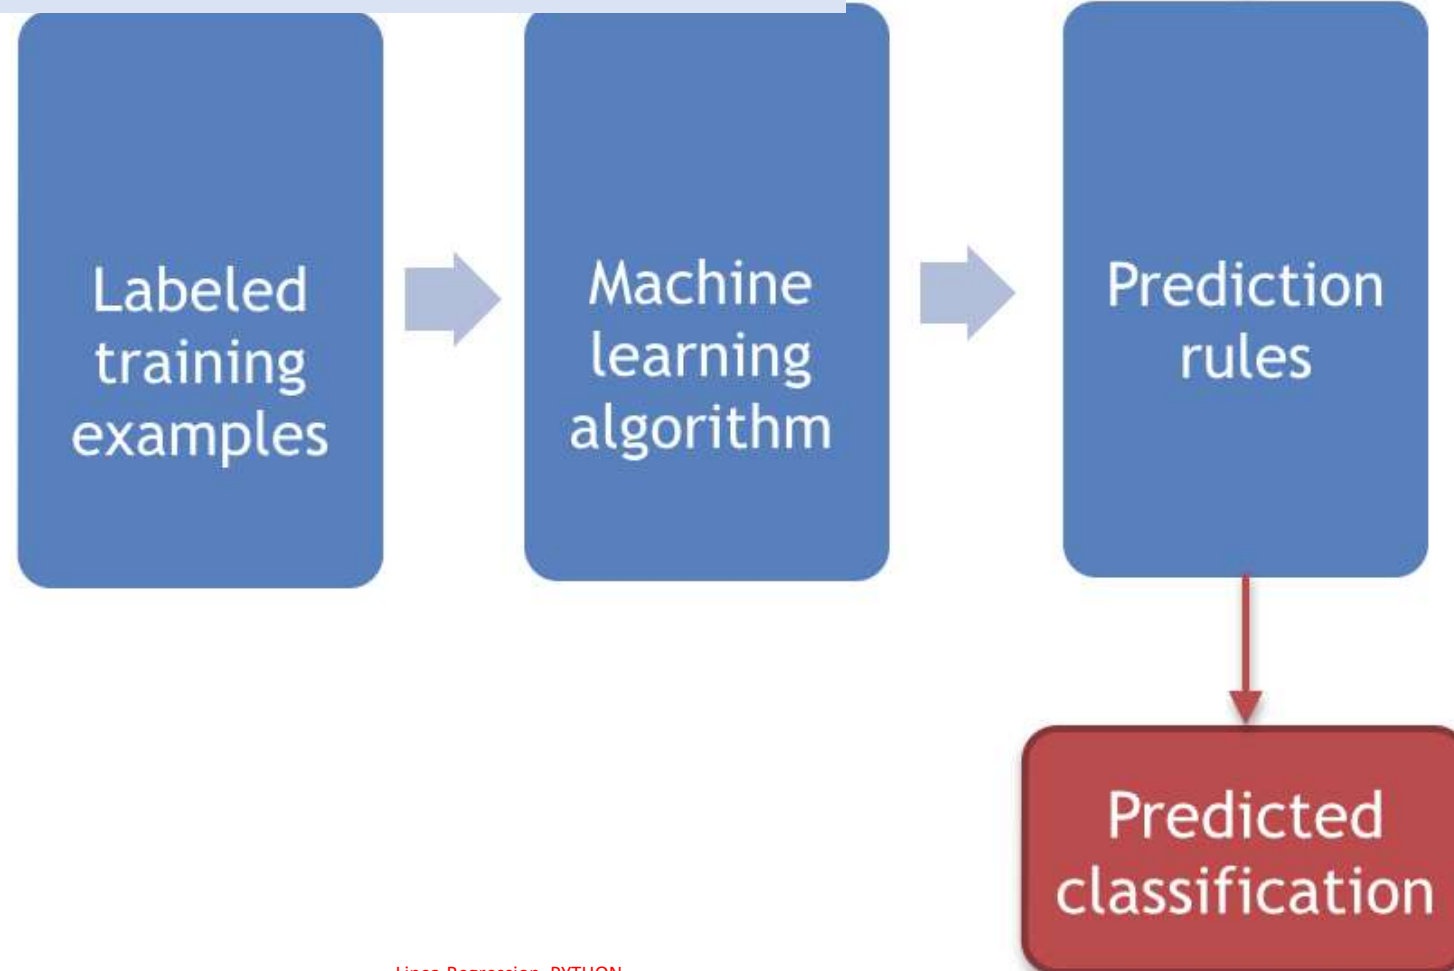
Regression

Output data  
continuous

Unsupervised Learning  
Group and analyze data  
based only on input data

Clustering

**Supervised Learning:** Applications in which the training data comprises examples of the input vectors along with their corresponding target vectors are known as supervised learning problems



# Introduction



- Regression problems are supervised learning problems in which the response is continuous
- Linear regression is a technique that is useful for regression problems.
- Classification problems are supervised learning problems in which the response is categorical
- Benefits of linear regression
  - widely used
  - runs fast
  - easy to use (not a lot of tuning required)
  - highly interpretable
  - basis for many other methods

# Libraries



- Statsmodels
- scikit-learn

```
# imports
import pandas as pd
import seaborn as sns
import statsmodels.formula.api as smf
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn.cross_validation import
train_test_split
import numpy as np

# allow plots to appear directly in the notebook
%matplotlib inline
```



```
# read data into a DataFrame
```

```
data = pd.read_csv('advertising.csv', index_col=0)  
data.head()
```

```
# shape of the DataFrame  
data.shape
```

```
# visualize the relationship between the  
features and the response using scatterplots
```

```
sns.pairplot(data,  
x_vars=['TV','Radio','Newspaper'],  
y_vars='Sales', size=7, aspect=0.7)
```

```
sns.pairplot(data)
```

```
sns.pairplot(data.dropna())
```

```
sns.pairplot(data, diag_kind="kde")
```

```
sns.pairplot(data, diag_kind="kde", markers="+",
              plot_kws=dict(s=50, edgecolor="b", linewidth=1),
              diag_kws=dict(shade=True))
```

```
sns.pairplot(data,
              x_vars=["TV", "radio", "newspaper"],
              y_vars=["sales"], size=7, aspect=0.7)
```

```
sns.pairplot(data,
              x_vars=["TV", "radio"],
              y_vars=["newspaper", "sales"])
```

# Simple Linear Regression

- Simple linear regression is an approach for predicting a quantitative response using a single feature (or "predictor" or "input variable")
- It takes the following form:
- $y = \beta_0 + \beta_1 x$
- What does each term represent?
- $y$  is the response
- $x$  is the feature
- $\beta_0$  is the intercept
- $\beta_1$  is the coefficient for  $x$
- $\beta_0$  and  $\beta_1$  are called the model coefficients
- To create your model, you must "learn" the values of these coefficients. Once we've learned these coefficients, we can use the model to predict Sales.



# Estimating ("Learning") Model Coefficients



- Coefficients are estimated using the least squares criterion
- In other words, we find the line (mathematically) which minimizes the sum of squared residuals (or "sum of squared errors"):

What elements are present in the diagram?

The black dots are the observed values of  $x$  and  $y$

The blue line is our least squares line

The red lines are the residuals, which are the distances between the observed values and the least squares line

How do the model coefficients relate to the least squares line?

$\beta_0$  is the intercept (the value of  $y$  when  $x=0$ )

$\beta_1$  is the slope (the change in  $y$  divided by change in  $x$ )

Here is a graphical depiction of those calculations:

# EXAMPLE

Boston House Prices dataset



- =====
- Data Set Characteristics:
- :Number of Instances: 506
- :Number of Attributes: 13 numeric/categorical predictive
- :Median Value (attribute 14) is usually the target
- :Attribute Information (in order):
- - CRIM per capita crime rate by town
- - ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- - INDUS proportion of non-retail business acres per town
- - CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- - NOX nitric oxides concentration (parts per 10 million)
- - RM average number of rooms per dwelling
- - AGE proportion of owner-occupied units built prior to 1940
- - DIS weighted distances to five Boston employment centres
- - RAD index of accessibility to radial highways
- - TAX full-value property-tax rate per \$10,000
- - PTRATIO pupil-teacher ratio by town
- - B  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of blacks by town
- - LSTAT % lower status of the population
- - MEDV Median value of owner-occupied homes in \$1000's
- :Missing Attribute Values: None
- :Creator: Harrison, D. and Rubinfeld, D.L.

Linear Regression PYTHON

```
In [1]: import statsmodels.api as sm

In [2]: from sklearn import datasets ## imports datasets from scikit-learn
data = datasets.load_boston() ## Loads Boston dataset from datasets Library

In [3]: import numpy as np
import pandas as pd
# define the data/predictors as the pre-set feature names
df = pd.DataFrame(data.data, columns=data.feature_names)

# Put the target (housing value -- MEDV) in another DataFrame
target = pd.DataFrame(data.target, columns=["MEDV"])
```

```
In [4]: ## Without a constant

import statsmodels.api as sm

X = df["RM"]
y = target["MEDV"]

# Note the difference in argument order
model = sm.OLS(y, X).fit()
predictions = model.predict(X) # make the predictions by the model

# Print out the statistics
model.summary()
```

Out[4]:

Dep. Variable:	MEDV	R-squared:	0.901
Model:	OLS	Adj. R-squared:	0.901
Method:	Least Squares	F-statistic:	4615.
Date:	Sun, 02 Sep 2018	Prob (F-statistic):	3.74e-258
Time:	21:29:59	Log-Likelihood:	-1747.1
No. Observations:	506	AIC:	3496.
Df Residuals:	505	BIC:	3500.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
RM	3.6534	0.054	67.930	0.000	3.548	3.759

Omnibus:	83.295	Durbin-Watson:	0.493
Prob(Omnibus):	0.000	Jarque-Bera (JB):	152.507
Skew:	0.955	Prob(JB):	7.65e-34
Kurtosis:	4.894	Cond. No.	1.00

data.feature\_names and data.target would print the column names of the independent variables and the dependent variable, respectively

RM      average number of rooms per dwelling

MEDV    Median value of owner-occupied homes in \$1000's  
house value/price data as a target variable and 13 other variables are set as predictors.

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



OLS Regression Results

Dep. Variable:	MEDV	R-squared:	0.901
Model:	OLS	Adj. R-squared:	0.901
Method:	Least Squares	F-statistic:	4615.
Date:	Sun, 02 Sep 2018	Prob (F-statistic):	3.74e-256
Time:	21:29:59	Log-Likelihood:	-1747.1
No. Observations:	506	AIC:	3496.
Df Residuals:	505	BIC:	3500.
Df Model:	1		
Covariance Type:	nonrobust		
	coef	std err	t P> t  [0.025 0.975]
RM	3.6534	0.054	67.930 0.000 3.548 3.759
Omnibus:	83.295	Durbin-Watson:	0.493
Prob(Omnibus):	0.000	Jarque-Bera (JB):	152.507
Skew:	0.955	Prob(JB):	7.65e-34
Kurtosis:	4.894	Cond. No.	1.00

- 1-what's the dependent variable and the model and the method.
- 2-OLS stands for **Ordinary Least Squares** and the method "Least Squares" means that we're trying to **fit** a regression line that would minimize the square of distance from the regression line
- 3- number of observations.
- 4-Df of residuals and models relates to the degrees of freedom — "the number of values in the final calculation of a statistic that are free to vary."
- 5-The coefficient of 3.634 means that as the RM variable increase the predicted value of MDEV increases by 3.634.
- 6-the R-squared — the percentage of variance the model explains,
- 7-the standard error (is the standard deviation of the sampling distribution of a statistic, most commonly of the mean);
- 8- The t scores and p-values, for hypothesis test — the RM has statistically significant p-value; there is a 95% confidence intervals for the RM (meaning we predict at a 95% percent confidence that the value of RM is between 3.548 to 3.759).



```
import statsmodels.api as sm # import statsmodels
```

```
X = df["RM"] ## X usually means our input variables (or independent variables)
y = target["MEDV"] ## Y usually means our output/dependent variable
X = sm.add_constant(X) ## Let's add an intercept (beta_0) to our model
```

```
# Note the difference in argument order
model = sm.OLS(y, X).fit() ## sm.OLS(output, input)
predictions = model.predict(X)
```

```
# Print out the statistics
model.summary()
```

#### OLS Regression Results

Dep. Variable:	MEDV	R-squared:	0.484
Model:	OLS	Adj. R-squared:	0.483
Method:	Least Squares	F-statistic:	471.8
Date:	Sun, 02 Sep 2018	Prob (F-statistic):	2.49e-74
Time:	21:42:24	Log-Likelihood:	-1673.1
No. Observations:	506	AIC:	3350.
Df Residuals:	504	BIC:	3359.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-34.6706	2.650	-13.084	0.000	-39.877	-29.465
RM	9.1021	0.419	21.722	0.000	8.279	9.925

Omnibus:	102.585	Durbin-Watson:	0.684
Prob(Omnibus):	0.000	Jarque-Bera (JB):	612.449
Skew:	0.726	Prob(JB):	1.02e-133

Kurtosis:	8.190	Cond. No.	58.4
-----------	-------	-----------	------

Linea Regression PYTHON



# OLS Regression Results

Dep. Variable:	MEDV	R-squared:	0.484
Model:	OLS	Adj. R-squared:	0.483
Method:	Least Squares	F-statistic:	471.8
Date:	Sun, 02 Sep 2018	Prob (F-statistic):	2.49e-74
Time:	21:42:24	Log-Likelihood:	-1673.1
No. Observations:	506	AIC:	3350.
Df Residuals:	504	BIC:	3359.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-34.6706	2.650	-13.084	0.000	-39.877	-29.465
RM	9.1021	0.419	21.722	0.000	8.279	9.925

Omnibus:	102.585	Durbin-Watson:	0.684
Prob(Omnibus):	0.000	Jarque-Bera (JB):	612.449
Skew:	0.726	Prob(JB):	1.02e-133
Kurtosis:	8.190	Cond. No.	58.4

Interpreting the Table — **With the constant term the coefficients are different.** Without a constant we are forcing our model to go through the origin, but now we have a y-intercept at -34.67. We also changed the slope of the RM predictor from **3.634 to 9.1021**.

# fitting a regression model with more than one variable — ( using RM and LSTAT)



## Interpreting the Output

```
X = df[["RM", "LSTAT"]]
y = target["MEDV"]
model = sm.OLS(y, X).fit()
predictions = model.predict(X)
model.summary()
```

### OLS Regression Results

Dep. Variable:	MEDV	R-squared:	0.948
Model:	OLS	Adj. R-squared:	0.948
Method:	Least Squares	F-statistic:	4637.
Date:	Sun, 02 Sep 2018	Prob (F-statistic):	0.00
Time:	21:46:45	Log-Likelihood:	-1582.9
No. Observations:	506	AIC:	3170.
Df Residuals:	504	BIC:	3178.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
RM	4.9069	0.070	69.906	0.000	4.769	5.045
LSTAT	-0.6557	0.031	-21.458	0.000	-0.716	-0.596

Omnibus:	145.153	Durbin-Watson:	0.834
Prob(Omnibus):	0.000	Jarque-Bera (JB):	442.157
Skew:	1.351	Prob(JB):	9.70e-97
Kurtosis:	6.698	Cond. No.	4.72

1-higher R-squared value — 0.948, meaning that this model explains 94.8% of the variance in our dependent variable. Whenever we add variables to a regression model,  $R^2$  will be higher.

2- Both RM and LSTAT are statistically significant in predicting (or estimating) the median house value;

3- as RM increases by 1, MEDV will increase by 4.9069 and when LSTAT increases by 1, MEDV will decrease by -0.6557. (LSTAT is the percentage of lower status of the population )

# Linear Regression in SKLearn



```
from sklearn import linear_model
```

```
from sklearn import datasets ## imports datasets from scikit-learn  
data = datasets.load_boston() ## loads Boston dataset from datasets library
```

```
# define the data/predictors as the pre-set feature names  
df = pd.DataFrame(data.data, columns=data.feature_names)  
  
# Put the target (housing value -- MEDV) in another DataFrame  
target = pd.DataFrame(data.target, columns=["MEDV"])
```

```
X = df  
y = target["MEDV"]
```

```
lm = linear_model.LinearRegression()  
model = lm.fit(X,y)
```

```
predictions = lm.predict(X)  
print(predictions)[0:5]
```

```
[30.00821269 25.0298606  30.5702317  28.60814055 27.94288232 25.25940048  
-----
```



The  $R^2$  score of the model is 0.740(the percentage of explained variance of the predictions)



```
lm.score(X,y)
```

```
0.7406077428649428
```

```
lm.coef_
```

```
array([-1.07170557e-01,  4.63952195e-02,  2.08602395e-02,  2.68856140e+00,  
       -1.77957587e+01,  3.80475246e+00,  7.51061703e-04, -1.47575880e+00,  
        3.05655038e-01, -1.23293463e-02, -9.53463555e-01,  9.39251272e-03,  
       -5.25466633e-01])
```

```
lm.intercept_
```

```
36.4911032803614
```



```
import pandas as pd
import numpy as np
import itertools
from itertools import chain, combinations
import statsmodels.formula.api as smf
import scipy.stats as scipystats
import statsmodels.api as sm
import statsmodels.stats.stattools as stools
import statsmodels.stats as stats
from statsmodels.graphics.regressionplots import *
import matplotlib.pyplot as plt
import seaborn as sns
import copy
from sklearn.cross_validation import train_test_split
import math
import time
```

```
elemapi = pd.read_csv('elemapi.csv')
```

```
print (elemapi[['api00', 'acs_k3', 'meals', 'full']].head())
```

	api00	acs_k3	meals	full
0	693	16.0	67.0	76.0
1	570	15.0	92.0	79.0
2	546	17.0	97.0	68.0
3	571	20.0	90.0	87.0
4	478	18.0	89.0	87.0

Let's dive right in and perform a regression analysis using the variables api00, acs\_k3, meals and full.

These measure:

- 1- the academic performance of the school (api00)
- 2- the average class size in kindergarten through 3rd grade (acs\_k3)
- 3- the percentage of students receiving free meals (meals) - which is an indicator of poverty, and the percentage of teachers who have full teaching credentials (full).

We expect that better academic performance would be associated with lower class size, fewer students receiving free meals, and a higher percentage of teachers having full teaching credentials.



## 1-linear regression using formulas

```
print ('-'*40 + ' smf.ols in R formula ' + '-'*40 + '\n')
lm = smf.ols(formula = 'api00 ~ ell', data = elemapi).fit()
print( lm.summary())
plt.figure()
plt.scatter(elemapi.ell, elemapi.api00, c = 'g')
plt.plot(elemapi.ell, lm.params[0] + lm.params[1] * elemapi.ell)
plt.xlabel('ell')
plt.ylabel('api00')
plt.title("Linear Regression Plot")

print (elemapi[['api00', 'acs_k3', 'meals', 'full']].head())
```



```
Dep. Variable:    api00    R-squared:    0.589
Model:            OLS      Adj. R-squared:    0.588
Method:            Least Squares    F-statistic:    571.0
Date:            Mon, 03 Sep 2018    Prob (F-statistic):    6.54e-79
Time:            12:04:50    Log-Likelihood:    -2372.1
No. Observations:    400    AIC:    4748.
Df Residuals:    398    BIC:    4756.
Df Model:    1
Covariance Type:    nonrobust
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	785.8903	7.370	106.638	0.000	771.402	800.379
ell	-4.3961	0.184	-23.895	0.000	-4.758	-4.034

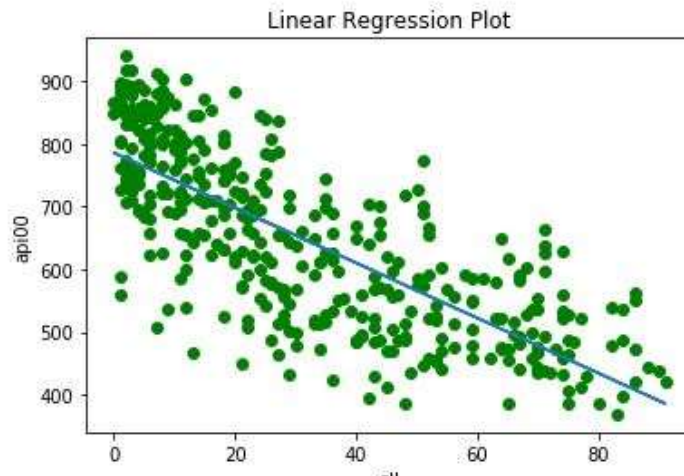
```
Omnibus:    7.668    Durbin-Watson:    1.461
Prob(Omnibus):    0.022    Jarque-Bera (JB):    7.264
Skew:    -0.283    Prob(JB):    0.0265
Kurtosis:    2.660    Cond. No.    64.7
```

#### Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

	api00	acs_k3	meals	full
0	693	16.0	67.0	76.0
1	570	15.0	92.0	79.0
2	546	17.0	97.0	68.0
3	571	20.0	90.0	87.0
4	478	18.0	89.0	87.0

## 1-linear regression using formulas



## 2-Simple linear regression using input directly



```
print ('-'*40 + ' sm.OLS with direct input data ' + '-'*40 + '\n')
lm2 = sm.OLS(elemapi['api00'], sm.add_constant(elemapi[['ell']])).fit()
print (lm2.summary())
```

----- sm.OLS with direct input data -----

### OLS Regression Results

```
=====
Dep. Variable:          api00    R-squared:                0.589
Model:                  OLS      Adj. R-squared:           0.588
Method:                 Least Squares    F-statistic:         571.0
Date:                   Mon, 03 Sep 2018    Prob (F-statistic):    6.54e-79
Time:                   12:08:58    Log-Likelihood:       -2372.1
No. Observations:       400    AIC:                  4748.
Df Residuals:           398    BIC:                  4756.
Df Model:                1
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	785.8903	7.370	106.638	0.000	771.402	800.379
ell	-4.3961	0.184	-23.895	0.000	-4.758	-4.034

```
=====
Omnibus:                 7.668    Durbin-Watson:           1.461
Prob(Omnibus):           0.022    Jarque-Bera (JB):        7.264
Skew:                    -0.283    Prob(JB):                0.0265
Kurtosis:                 2.660    Cond. No.                 64.7
=====
```



# 1-Multiple linear regression using formulas



```
lm = smf.ols(formula = 'api00 ~ acs_k3 + meals + full', data = elemapi).fit()
print (lm.summary())
```

## OLS Regression Results

```
=====
Dep. Variable:          api00    R-squared:                0.674
Model:                  OLS      Adj. R-squared:           0.671
Method:                 Least Squares    F-statistic:         213.4
Date:                  Mon, 03 Sep 2018    Prob (F-statistic):    5.73e-75
Time:                  12:12:57    Log-Likelihood:       -1744.6
No. Observations:      313        AIC:                  3497.
Df Residuals:          309        BIC:                  3512.
Df Model:               3
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	906.7392	28.265	32.080	0.000	851.123	962.355
acs_k3	-2.6815	1.394	-1.924	0.055	-5.424	0.061
meals	-3.7024	0.154	-24.038	0.000	-4.005	-3.399
full	0.1086	0.091	1.197	0.232	-0.070	0.287

```
=====
Omnibus:                2.012    Durbin-Watson:           1.467
Prob(Omnibus):           0.366    Jarque-Bera (JB):        2.070
Skew:                    0.162    Prob(JB):                0.355
Kurtosis:                2.767    Cond. No.                 769.
=====
```

# 1-Multiple linear regression using input (drop missing values)

```
data = elemapi[['api00', 'acs_k3', 'meals', 'full']]
data = data.dropna(axis = 0, how = 'any')

lm2 = sm.OLS(data['api00'], sm.add_constant(data[['acs_k3', 'meals', 'full']])).fit()
print( lm2.summary())
```

## OLS Regression Results

```
=====
Dep. Variable:          api00    R-squared:                0.674
Model:                  OLS      Adj. R-squared:           0.671
Method:                 Least Squares    F-statistic:          213.4
Date:                  Mon, 03 Sep 2018    Prob (F-statistic):    5.73e-75
Time:                  12:19:50    Log-Likelihood:        -1744.6
No. Observations:      313    AIC:                  3497.
Df Residuals:          309    BIC:                  3512.
Df Model:               3
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	906.7392	28.265	32.080	0.000	851.123	962.355
acs_k3	-2.6815	1.394	-1.924	0.055	-5.424	0.061
meals	-3.7024	0.154	-24.038	0.000	-4.005	-3.399
full	0.1086	0.091	1.197	0.232	-0.070	0.287

```
=====
Omnibus:                2.012    Durbin-Watson:          1.467
Prob(Omnibus):           0.366    Jarque-Bera (JB):        2.070
Skew:                   0.162    Prob(JB):                0.355
Kurtosis:               2.767    Cond. No.                 769.
=====
```

# Data Analysis



## Numeric Data Analysis

1. if there is any missing data
2. what is the distribution of the data, and its visualization like histogram and box-plot
3. what is the five numbers: min, 25 percentile, median, 75 percentile, and the max
4. mean, stdev, length
5. the correlation between the data
6. furthermore, is there any outliers in the data?
7. other plots: pairwise scatter plot, kernel density plot

## Categorical Data Analysis

1. is there any missing data?
2. how many unique values of the data? what is their frequency?

Feature selection !?



```
sample_data = elemapi(['api00', 'acs_k3', 'meals', 'full'])
print (sample_data.describe())
```

	api00	acs_k3	meals	full
count	400.000000	398.000000	315.000000	400.000000
mean	647.622500	18.547739	71.993651	66.056800
std	142.248961	5.004933	24.385570	40.297926
min	369.000000	-21.000000	6.000000	0.420000
25%	523.750000	18.000000	57.000000	0.950000
50%	643.000000	19.000000	77.000000	87.000000
75%	762.250000	20.000000	93.000000	97.000000
max	940.000000	25.000000	100.000000	100.000000

1. api00 and full does not have missing values and their length is 400.
2. acs\_k3 has 398 non-missing values, so it has two missing data. meals has 315 non-missing values so it has 85 missings

```
from scipy.stats import gaussian_kde
```

```
plt.hist(elemapi.api00, 20, normed = 1, facecolor = 'g', alpha = 0.5)
```

```
# add density plot
```

```
density = gaussian_kde(elemapi.api00)
```


```
xs = np.linspace(300, 1000, 500)
```

```
density.covariance_factor = lambda : .2
```

```
density._compute_covariance()
```


```
plt.plot(xs, density(xs), color = "b")
```

```
plt.show()
```



```
plt.boxplot(elemapi.api00, 0, 'gD')
```

```
{'whiskers': [<matplotlib.lines.Line2D at 0x1c070566dd8>,
<matplotlib.lines.Line2D at 0x1c07056f2b0>],
'caps': [<matplotlib.lines.Line2D at 0x1c07056f6d8>,
<matplotlib.lines.Line2D at 0x1c07056fb00>],
'boxes': [<matplotlib.lines.Line2D at 0x1c070566c88>],
'medians': [<matplotlib.lines.Line2D at 0x1c07056ff28>],
'fliers': [<matplotlib.lines.Line2D at 0x1c070576390>],
'means': []}
```





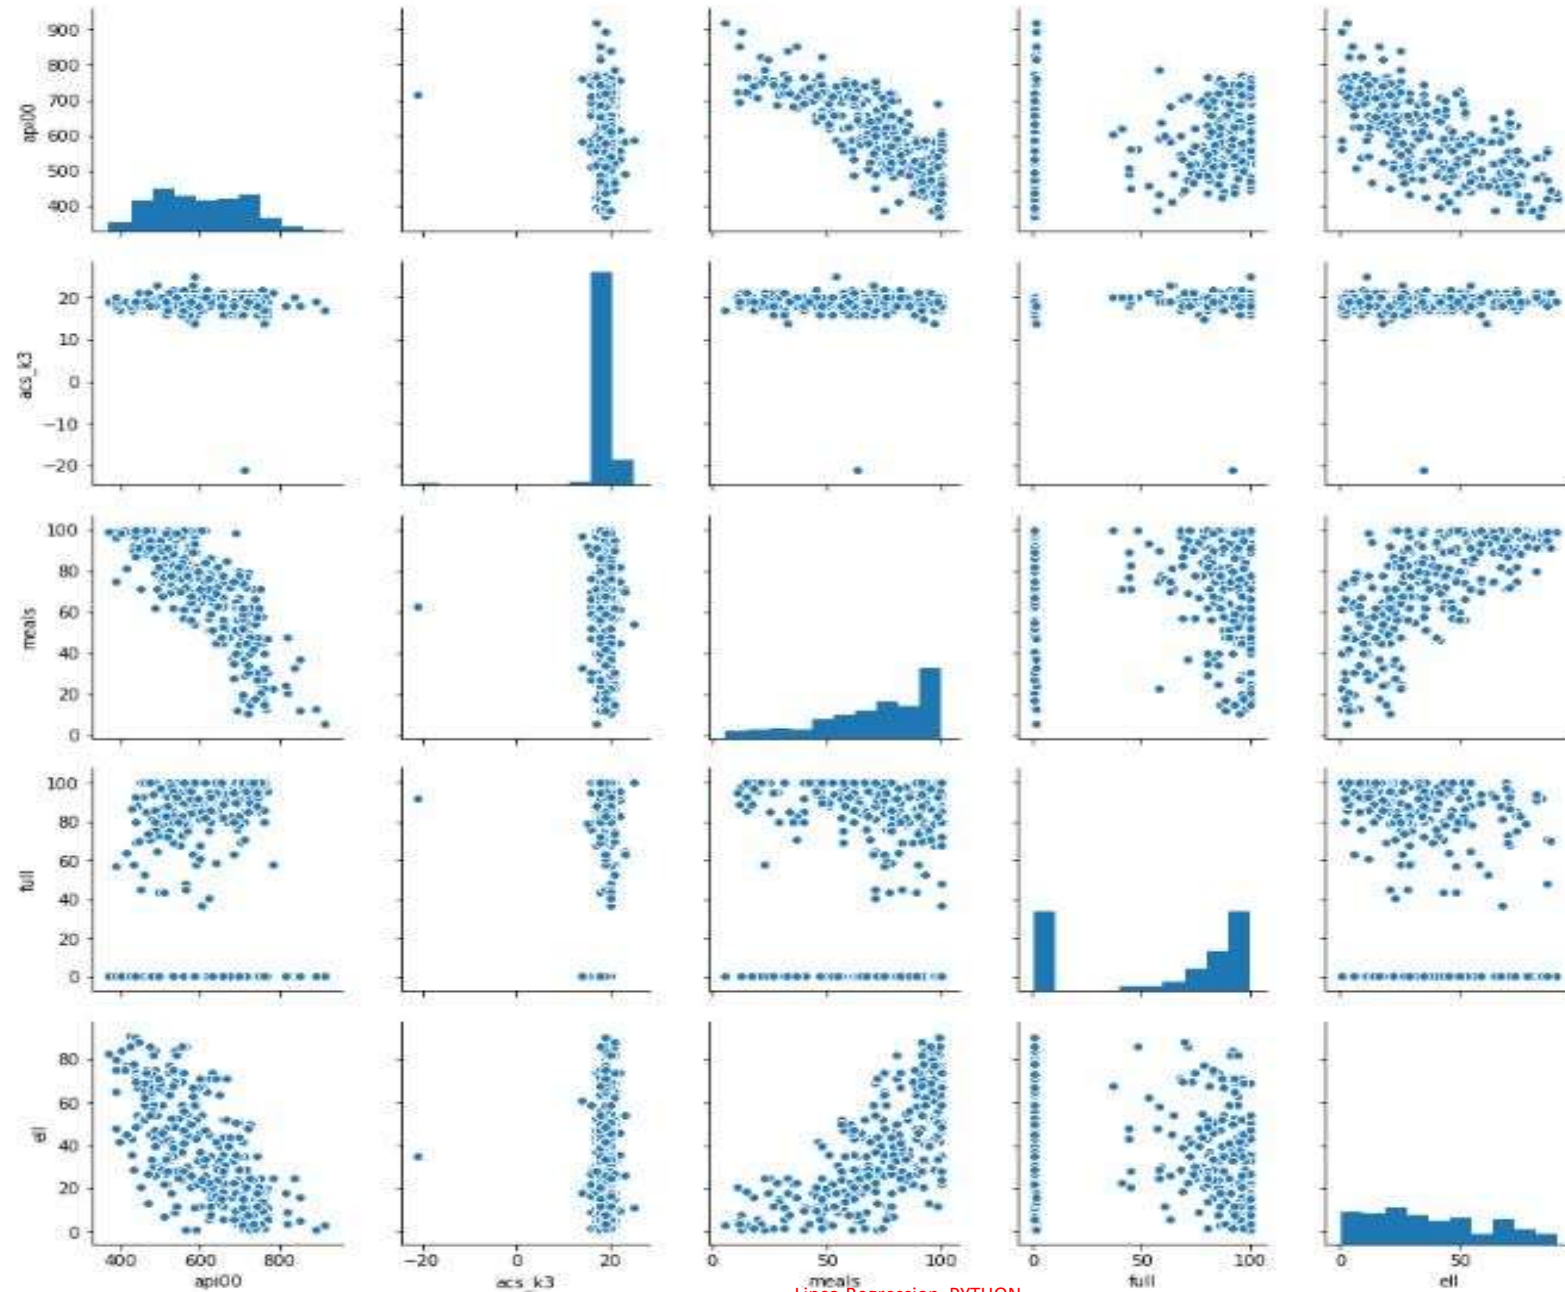
```
# check correlation between each variable and api00  
print (elemapi.corr().loc['api00', :].sort_values())
```

```
mealcat    -0.867260  
meals      -0.819300  
ell        -0.767634  
not_hsg    -0.683255  
emer       -0.582731  
yr_rnd     -0.475440  
hsg        -0.355809  
enroll     -0.318172  
mobility   -0.206410  
growth     -0.108158  
acs_k3     -0.095546  
dnum       -0.011383  
snum       0.216457  
acs_46     0.232912  
some_col   0.261527  
full       0.411125  
col_grad   0.527301  
grad_sch   0.633241  
avg_ed     0.792954  
api99      0.985343  
api00      1.000000  
Name: api00, dtype: float64
```



```
sns.pairplot(elemapi[['api00', 'acs_k3', 'meals', 'full', 'ell']].dropna(how = 'any', axis = 0))
```

<seaborn.axisgrid.PairGrid at 0x1c070584fd0>




# kernel density function => the ell variable skewed to the right.



```
sns.kdeplot(np.array(elemapi.ell), bw=0.5)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x1c0710cb390>
```



# Categorical variable frequency



```
elemapi.acs_k3.value_counts(dropna = False).sort_index()
```

```
-21.0      3
-20.0      2
-19.0      1
 14.0      2
 15.0      1
 16.0     14
 17.0     20
 18.0     64
 19.0    143
 20.0     97
 21.0     40
 22.0      7
 23.0      3
 25.0      1
NaN         2
Name: acs_k3, dtype: int64
```

for acs\_k3, there are about 143 data points equal to 19.

# Regression Diagnostics



verifying that the data have met the regression assumptions



## Issues that can arise during the analysis



- **Linearity** - the relationships between the predictors and the outcome variable should be linear
- **Normality** - the errors should be normally distributed - technically normality is necessary only for the t-tests to be valid, estimation of the coefficients only requires that the errors be identically and independently distributed
- **Homogeneity** of variance (homoscedasticity) - the error variance should be constant
- **Independence** - the errors associated with one observation are not correlated with the errors of any other observation
- **Errors in variables** - predictor variables are measured without error
- **Model specification** - the model should be properly specified (including all relevant variables, and excluding irrelevant variables)

# Unusual and influential data



1. **Outliers:** In linear regression, an outlier is an observation with **large residual**. In other words, it is an observation whose dependent-variable value is unusual given its values on the predictor variables. An outlier may indicate a sample peculiarity or may indicate a data entry error or other problem.
2. **Leverage:** An observation with an extreme value on a predictor variable is called a point with high leverage. **Leverage is a measure of how far an observation deviates from the mean of that variable.** These leverage points can have an effect on the estimate of regression coefficients.
3. **Influence:** An observation is said to be influential if removing the observation substantially changes the estimate of coefficients. Influence can be thought of as the **product of leverage and outlierness.**



## Data set crime

```
import pandas as pd
import numpy as np
import itertools
from itertools import chain, combinations
import statsmodels.formula.api as smf
import scipy.stats as scipystats
import statsmodels.api as sm
import statsmodels.stats.stattools as stools
import statsmodels.stats as stats
from statsmodels.graphics.regressionplots import *
import matplotlib.pyplot as plt
import seaborn as sns
import copy
from sklearn.cross_validation import train_test_split
import math
import time
```

```
%matplotlib inline
plt.rcParams['figure.figsize'] = (16, 12)
```

```
crime = pd.read_csv('crime.csv')
```

```
crime.describe()
```

	low	murder	tc2009
count	48.000000	48.000000	48.000000
mean	-40.104167	4.495833	3388.716667
std	17.786248	2.354530	749.441028
min	-80.000000	0.900000	2026.200000
25%	-50.250000	2.675000	2778.825000
50%	-40.000000	4.650000	3442.850000
75%	-29.750000	5.925000	4018.575000
max	12.000000	12.300000	4562.200000





# Advanced

# Regression with Categorical Predictors



## Regression with categorical predictors

- 1 Regression with a 0/1 variable
- 2 Regression with a 1/2 variable
- 3 Regression with a 1/2/3 variable
- 4 Regression with multiple categorical predictors
- 5 Categorical predictor with interactions
- 6 Continuous and categorical variables
- 7 Interactions of continuous by 0/1 categorical variables
- 8 Continuous and categorical variables, interaction with 1/2/3 variable

# Regression with a 0/1 variable

```
import warnings
```

```
with warnings.catch_warnings():  
    warnings.filterwarnings("ignore", category=DeprecationWarning)
```

```
import pandas as pd  
import numpy as np  
import itertools  
from itertools import chain, combinations  
import statsmodels.formula.api as smf  
import scipy.stats as scipystats  
import statsmodels.api as sm  
import statsmodels.stats.stattools as stools  
import statsmodels.stats as stats  
from statsmodels.graphics.regressionplots import *  
import matplotlib.pyplot as plt  
import seaborn as sns  
import copy  
from sklearn.cross_validation import train_test_split  
import math  
import time
```

```
%matplotlib inline
```

```
elemapi2 = pd.read_csv('elemapi.csv')
```

```
elemapi2_sel = elemapi2.loc[:, ["api00", "some_col", "yr_rnd", "mealcat"]]
```







```
print(elemapi2_sel.describe())

def cv_desc(df, var):
    return df[var].value_counts(dropna = False)
```

```
print ('\n' )
print( cv_desc(elemapi2_sel, 'mealcat'))
print ('\n' )
print (cv_desc(elemapi2_sel, 'yr_rnd'))
```

	api00	some_col	yr_rnd	mealcat
count	400.000000	400.000000	400.000000	400.000000
mean	647.622500	19.712500	0.230000	2.015000
std	142.248961	11.336938	0.421360	0.819423
min	369.000000	0.000000	0.000000	1.000000
25%	523.750000	12.000000	0.000000	1.000000
50%	643.000000	19.000000	0.000000	2.000000
75%	762.250000	28.000000	0.000000	3.000000
max	940.000000	67.000000	1.000000	3.000000

```
3    137
2    132
1    131
Name: mealcat, dtype: int64
```

```
0    308
1     92
Name: yr_rnd, dtype: int64
```

1. The variable `api00` is a measure of the performance of the students.
2. The variable `some_col` is a continuous variable that measures the percentage of the parents in the school who have attended college.
3. The variable `yr_rnd` is a categorical variable that is coded 0 if the school is not year round, and 1 if year round.
4. The variable `meals` is the percentage of students who are receiving state sponsored free meals and can be used as an indicator of poverty. This was broken into 3 categories (to make equally sized groups) creating the variable `mealcat`. e.



Macro function gives codebook type information on a specific variable.



```
def codebook(df, var):
    title = "Codebook for " + str(var)
    unique_values = len(df[var].unique())
    max_v = df[var].max()
    min_v = df[var].min()
    n_miss = sum(pd.isnull(df[var]))
    mean = df[var].mean()
    stdev = df[var].std()
    print(pd.DataFrame({'title': title, 'unique values': unique_values, 'max value' : max_v, 'min value' : min_v, 'num of missing' : n_miss, 'mean' : mean, 'stdev' : stdev}))
    return

codebook(elemapi2_sel, 'api00')
```

```

      title  unique values  max value  min value  num of missing  \
0  Codebook for api00      271      940      369              0

      mean      stdev
0  647.6225  142.248961
```



```
reg = smf.ols(formula = "api00 ~ yr_rnd", data = elemapi2_sel).fit()  
reg.summary()
```

#### OLS Regression Results

Dep. Variable:	api00	R-squared:	0.226
Model:	OLS	Adj. R-squared:	0.224
Method:	Least Squares	F-statistic:	116.2
Date:	Mon, 03 Sep 2018	Prob (F-statistic):	5.96e-24
Time:	09:19:25	Log-Likelihood:	-2498.9
No. Observations:	400	AIC:	5002.
Df Residuals:	398	BIC:	5010.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	684.5390	7.140	95.878	0.000	670.503	698.575
yr_rnd	-160.5064	14.887	-10.782	0.000	-189.774	-131.239

Omnibus:	45.748	Durbin-Watson:	1.499
Prob(Omnibus):	0.000	Jarque-Bera (JB):	13.162
Skew:	0.006	Prob(JB):	0.00139
Kurtosis:	2.111	Cond. No.	2.53

## Regression with a 0/1 variable (categorical dummy variable)

$$\text{api00} = \text{Intercept} + \text{Byr\_rnd} * \text{yr\_rnd}$$

$$\text{api00} = 684.539 + -160.5064 * \text{yr\_rnd}$$

yr\_rnd is 0

$$\text{api00} = \text{constant} + 0 * \text{Byr\_rnd}$$

$$\begin{aligned}\text{api00} &= 684.539 + 0 * -160.5064 \\ \text{api00} &= 684.539\end{aligned}$$

yr\_rnd is 1

$$\text{api00} = \text{constant} + 1 * \text{Byr\_rnd}$$

$$\text{api00} = 684.539 + 1 * -160.5064$$


$$\text{api00} = 524.0326$$

predicted value for the year round  
schools is 524.032

# Predicted values to the mean api00 scores for the year-round and non-year-round students



```
plt.scatter(elemapi2_sel.yr_rnd, elemapi2_sel.api00, marker = "+")
plt.plot([0, 1], [np.mean(elemapi2_sel.query('yr_rnd == 0').api00), np.mean(elemapi2_sel.query('yr_rnd == 1').api00)], 'r--')
plt.ylabel("api 2000")
plt.xlabel("year round school")
plt.show()
```



```
elemapi2_sel["yr_rnd_c"] = elemapi2_sel.yr_rnd.map({0: "No", 1: "Yes"})
elemapi2_sel["mealcat_c"] = elemapi2_sel.mealcat.map({1: "0-46% free meals", 2: "47-80% free meals", 3: "1-100% free meals"})

elemapi2_sel_group = elemapi2_sel.groupby("yr_rnd_c")
elemapi2_sel_group.api00.agg([np.mean, np.std])
```

	mean	std
yr_rnd_c		
No	684.538961	132.112534
Yes	524.032609	98.916043

For the non-year-round schools, their mean is the same as the intercept (684.539). The coefficient for yr\_rnd is the amount we need to add to get the mean for the year-round schools, i.e., we need to add -160.5064 to get 524.0326, the mean for the non year-round schools.

→ Byr\_rnd is the mean api00 score for the year-round schools minus the mean api00 score for the non year-round schools, i.e.,  $\text{mean}(\text{year-round}) - \text{mean}(\text{non year-round})$ .



The t value below is the same as the t value for yr\_rnd in the regression.

```
# pooled ttest, assume equal population variance
print( scipystats.ttest_ind(elemapi2_sel.query('yr_rnd == 0').api00,
elemapi2_sel.query('yr_rnd == 1').api00))
```

```
# does not assume equal variance
print (scipystats.ttest_ind(elemapi2_sel.query('yr_rnd == 0').api00,
elemapi2_sel.query('yr_rnd == 1').api00, equal_var = False))
```

Ttest\_indResult(statistic=10.781500136400451, pvalue=5.9647081127888056e-24)

Ttest\_indResult(statistic=12.57105956566846, pvalue=5.29731480664924e-27)

Since a t-test is the same as doing an ANOVA (analysis of variance:test whether a number of me

```
print (scipystats.f_oneway(elemapi2_sel.query('yr_rnd == 0').api00, elemapi2_sel.query('yr_rnd
== 1').api00))
```

F\_onewayResult(statistic=116.2407451912029, pvalue=5.964708112790799e-24)

If we square the t-value from the t-test, we get the same value as the F-value from ANOVA:  $10.78^2=116.21$  (with a little rounding error.)



```
elemapi2_sel["yr_rnd2"] = elemapi2_sel["yr_rnd"] + 1
reg = smf.ols("api00 ~ yr_rnd2", data = elemapi2_sel).fit()
reg.summary()
```

#### OLS Regression Results

<b>Dep. Variable:</b>	api00	<b>R-squared:</b>	0.226			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.224			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	116.2			
<b>Date:</b>	Mon, 03 Sep 2018	<b>Prob (F-statistic):</b>	5.96e-24			
<b>Time:</b>	10:02:17	<b>Log-Likelihood:</b>	-2498.9			
<b>No. Observations:</b>	400	<b>AIC:</b>	5002.			
<b>Df Residuals:</b>	398	<b>BIC:</b>	5010.			
<b>Df Model:</b>	1					
<b>Covariance Type:</b>	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	845.0453	19.353	43.664	0.000	806.998	883.093
yr_rnd2	-160.5064	14.887	-10.782	0.000	-189.774	-131.239
<b>Omnibus:</b>	45.748	<b>Durbin-Watson:</b>	1.499			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	13.162			
<b>Skew:</b>	0.006	<b>Prob(JB):</b>	0.00139			
<b>Kurtosis:</b>	2.111	<b>Cond. No.</b>	6.23			

Linea Regressione

## Regression with a 1/2 variable



a copy of the variable yr\_rnd called yr\_rnd2 that is coded 1/2, 1=non year-round and 2=year-round.

Note that the coefficient for yr\_rnd is the same as yr\_rnd2. So, you can see that if you code yr\_rnd as 0/1 or as 1/2, the regression coefficient works out to be the same. However the intercept (Intercept) is a bit less intuitive. When we used yr\_rnd, the intercept was the mean for the non year-rounds. When using yr\_rnd2, the intercept is the mean for the non year-rounds minus Byr\_rnd2, i.e.,  $684.539 - (-160.506) = 845.045$



# The relationship between the amount of poverty and api scores?

```
elemapi2_sel_group = elemapi2_sel.groupby("mealcat_c")
elemapi2_sel_group.api00.agg([lambda x: x.shape[0], np.mean, np.std])
```

	<lambda>	mean	std
mealcat_c			
0-46% free meals	131	805.717557	65.668664
1-100% free meals	137	504.379562	62.727015
47-80% free meals	132	639.393939	82.135130

Regression with a 1/2/3 variable

use mealcat as a **proxy for a measure of poverty**. Mealcat ( has three unique values on free meals. We can associate a value label to variable mealcat to make it more meaningful for us when we run python regression with mealcat.

```
lm = smf.ols('api00 ~ mealcat', data = elemapi2_sel).fit()
lm.summary()
```

#### OLS Regression Results

<b>Dep. Variable:</b>	api00	<b>R-squared:</b>	0.752
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.752
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1208.
<b>Date:</b>	Mon, 03 Sep 2018	<b>Prob (F-statistic):</b>	1.29e-122
<b>Time:</b>	10:10:41	<b>Log-Likelihood:</b>	-2271.1
<b>No. Observations:</b>	400	<b>AIC:</b>	4546.
<b>Df Residuals:</b>	398	<b>BIC:</b>	4554.
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	950.9874	9.422	100.935	0.000	932.465	969.510
<b>mealcat</b>	-150.5533	4.332	-34.753	0.000	-159.070	-142.037

<b>Omnibus:</b>	3.106	<b>Durbin-Watson:</b>	1.516
<b>Prob(Omnibus):</b>	0.212	<b>Jarque-Bera (JB):</b>	3.112
<b>Skew:</b>	-0.214	<b>Prob(JB):</b>	0.211
<b>Kurtosis:</b>	2.943	<b>Cond. No.</b>	6.86



mealcat is not an interval variable. (need to create dummy variables). For example, in order to create dummy variables for mealcat, we can do the following using sklearn to create dummy variables.



```
from sklearn import preprocessing
le_mealcat = preprocessing.LabelEncoder()
elemapi2_sel['mealcat_dummy'] =
le_mealcat.fit_transform(elemapi2_sel.mealcat)
```

```
elemapi2_sel.groupby('mealcat_dummy').size()
```

```
ohe = preprocessing.OneHotEncoder()
dummy = pd.DataFrame(ohe.fit_transform(elemapi2_sel.mealcat.reshape(-
1,1)).toarray(), columns = ["mealcat1", "mealcat2", "mealcat3"])
elemapi2_sel = pd.concat([elemapi2_sel, dummy], axis = 1)
```

```
lm = smf.ols('api00 ~ mealcat2 + mealcat3', data = elemapi2_sel).fit()
lm.summary()
```

	coef	std err	t	P> t	[95.0% Conf. Int.]
<b>Intercept</b>	805.7176	6.169	130.599	0.000	793.589 817.846
<b>mealcat2[0]</b>	-83.1618	4.354	-19.099	0.000	-91.722 - 74.602
<b>mealcat2[1]</b>	-83.1618	4.354	-19.099	0.000	-91.722 - 74.602
<b>mealcat3[0]</b>	-150.6690	4.314	-34.922	0.000	-159.151 - 142.187
<b>mealcat3[1]</b>	-150.6690	4.314	-34.922	0.000	-159.151 - 142.187

```
lm = lm = smf.ols('api00 ~ C(mealcat)', data = elemapi2_sel).fit()
lm.summary()
```

#### OLS Regression Results

Dep. Variable:	api00	R-squared:	0.755
Model:	OLS	Adj. R-squared:	0.754
Method:	Least Squares	F-statistic:	611.1
Date:	Mon, 03 Sep 2018	Prob (F-statistic):	6.48e-122
Time:	10:25:35	Log-Likelihood:	-2269.0
No. Observations:	400	AIC:	4544.
Df Residuals:	397	BIC:	4556.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	805.7176	6.169	130.599	0.000	793.589	817.846
C(mealcat)[T.2]	-166.3236	8.708	-19.099	0.000	-183.444	-149.203
C(mealcat)[T.3]	-301.3380	8.629	-34.922	0.000	-318.302	-284.374

Omnibus:	1.593	Durbin-Watson:	1.541
Prob(Omnibus):	0.451	Jarque-Bera (JB):	1.684
Skew:	-0.139	Prob(JB):	0.431
Kurtosis:	2.847	Cond. No.	3.76

Run regression  
with categorical  
variable directly



## Regression with two categorical predictors



```
lm = smf.ols('api00 ~ yr_rnd', data = elemapi2_sel).fit()
print(lm.summary())

print('\n')

lm = smf.ols('api00 ~ mealcat1 + mealcat2', data = elemapi2_sel).fit()
print(lm.summary())
```

### OLS Regression Results

```
=====
Dep. Variable:          api00    R-squared:                0.226
Model:                  OLS      Adj. R-squared:           0.224
Method:                 Least Squares    F-statistic:          116.2
Date:                  Mon, 03 Sep 2018    Prob (F-statistic):    5.96e-24
Time:                  10:28:57    Log-Likelihood:        -2498.9
No. Observations:      400    AIC:                   5002.
Df Residuals:          398    BIC:                   5010.
Df Model:              1
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	684.5390	7.140	95.878	0.000	670.503	698.575
yr_rnd	-160.5064	14.887	-10.782	0.000	-189.774	-131.239

```
=====
Omnibus:                45.748    Durbin-Watson:           1.499
Prob(Omnibus):          0.000    Jarque-Bera (JB):        13.162
Skewness:               0.000    Prob(JB):                0.00130
Kurtosis:               3.000    Heteroskedasticity (H):  0.0000001
=====
```



```
lm = smf.ols('api00 ~ yr_rnd + C(mealcat)', data = elemapi2_sel).fit()
print(lm.summary())
```

### OLS Regression Results

```
=====
Dep. Variable:          api00    R-squared:                0.767
Model:                  OLS      Adj. R-squared:           0.765
Method:                 Least Squares    F-statistic:          435.0
Date:                   Mon, 03 Sep 2018    Prob (F-statistic):    6.40e-125
Time:                   10:34:47    Log-Likelihood:       -2258.6
No. Observations:       400    AIC:                  4525.
Df Residuals:           396    BIC:                  4541.
Df Model:                3
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	808.0131	6.040	133.777	0.000	796.139	819.888
C(mealcat)[T.2]	-163.7374	8.515	-19.229	0.000	-180.478	-146.997
C(mealcat)[T.3]	-281.6832	9.446	-29.821	0.000	-300.253	-263.113
yr_rnd	-42.9601	9.362	-4.589	0.000	-61.365	-24.555

```
=====
Omnibus:                1.210    Durbin-Watson:           1.584
Prob(Omnibus):           0.546    Jarque-Bera (JB):        1.279
Skew:                    -0.086    Prob(JB):                0.528
Kurtosis:                2.783    Cond. No.                 4.17
=====
```

[Control Panel]



## The effect of yr\_rnd at each of the three levels of mealcat.



```
lm = smf.ols('api00 ~ C(yr_rnd) * C(mealcat)', data = elemapi2_sel).fit()
print (lm.summary())
```

### OLS Regression Results

=====						
Dep. Variable:	api00	R-squared:	0.769			
Model:	OLS	Adj. R-squared:	0.766			
Method:	Least Squares	F-statistic:	261.6			
Date:	Mon, 03 Sep 2018	Prob (F-statistic):	9.19e-123			
Time:	10:45:35	Log-Likelihood:	-2257.5			
No. Observations:	400	AIC:	4527.			
Df Residuals:	394	BIC:	4551.			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
Intercept	809.6855	6.185	130.911	0.000	797.526	821.845
C(yr_rnd)[T.1]	-74.2569	26.756	-2.775	0.006	-126.860	-21.654
C(mealcat)[T.2]	-164.4120	8.877	-18.522	0.000	-181.864	-146.960
C(mealcat)[T.3]	-288.1929	10.443	-27.597	0.000	-308.724	-267.662
C(yr_rnd)[T.1]:C(mealcat)[T.2]	22.5167	32.752	0.687	0.492	-41.873	86.907
C(yr_rnd)[T.1]:C(mealcat)[T.3]	40.7644	29.231	1.395	0.164	-16.704	98.233
=====						
Omnibus:	1.439	Durbin-Watson:	1.583			
Prob(Omnibus):	0.487	Jarque-Bera (JB):	1.484			
Skew:	-0.096	Prob(JB):	0.476			
Kurtosis:	2.771	Cond. No.	16.4			
=====						

results show no indication of

results show no indication of interaction

```
elemapi2_ = elemapi2.copy()

lm = smf.ols('api00 ~ yr_rnd + some_col', data = elemapi2_).fit()
print(lm.summary())

elemapi2_['pred'] = lm.predict()


plt.scatter(elemapi2_.query('yr_rnd == 0').some_col, elemapi2_.query('yr_rnd == 0').pred, c = "b", marker = "o")
plt.scatter(elemapi2_.query('yr_rnd == 1').some_col, elemapi2_.query('yr_rnd == 1').pred, c = "r", marker = "+")
plt.plot()
```

## OLS Regression Results

```
=====
Dep. Variable:          api00    R-squared:                0.257
Model:                  OLS      Adj. R-squared:           0.253
Method:                 Least Squares    F-statistic:           68.54
Date:                  Mon, 03 Sep 2018    Prob (F-statistic):    2.69e-26
Time:                  10:49:35    Log-Likelihood:        -2490.8
No. Observations:      400    AIC:                   4988.
Df Residuals:          397    BIC:                   5000.
Df Model:               2
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	637.8581	13.503	47.237	0.000	611.311	664.405
yr_rnd	-149.1591	14.875	-10.027	0.000	-178.403	-119.915
some_col	2.2357	0.553	4.044	0.000	1.149	3.323

```
=====
Omnibus:                23.070    Durbin-Watson:           1.565
Prob(Omnibus):           0.000    Jarque-Bera (JB):        9.935
Skew:                    0.125    Prob(JB):                0.00696
Kurtosis:                 2.269    Cond. No.                 62.5
=====
```



Relationship between some\_col and api00 but there were two regression lines, one higher than the other but with equal slope. Such a model assumed that the slope was the same for the two groups. Perhaps the slope might be different!!

	coef
Intercept	637.8581
yr_rnd	-149.1591
some_col	2.2357

- The coefficient for some\_col indicates that for every unit increase in some\_col the api00 score is predicted to increase by 2.23 units.
- The graph has two lines, one for the year round schools and one for the non-year round schools.
- The coefficient for yr\_rnd is -149.16, indicating that as yr\_rnd increases by 1 unit, the api00 score is expected to decrease by about 149 units.
- the top line is about 150 units higher than the lower line.
- the intercept is 637 and that is where the upper line crosses the Y axis when X is 0. The lower line crosses the line about 150 units lower at about 487.



## Using categorical variable directly



```
lm = smf.ols('api00 ~ C(yr_rnd) + some_col', data = elemapi2_).fit()  
print (lm.summary())
```

### OLS Regression Results

```
=====
```

Dep. Variable:	api00	R-squared:	0.257
Model:	OLS	Adj. R-squared:	0.253
Method:	Least Squares	F-statistic:	68.54
Date:	Mon, 03 Sep 2018	Prob (F-statistic):	2.69e-26
Time:	10:57:12	Log-Likelihood:	-2490.8
No. Observations:	400	AIC:	4988.
Df Residuals:	397	BIC:	5000.
Df Model:	2		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
Intercept	637.8581	13.503	47.237	0.000	611.311	664.405
C(yr_rnd)[T.1]	-149.1591	14.875	-10.027	0.000	-178.403	-119.915
some_col	2.2357	0.553	4.044	0.000	1.149	3.323

```
=====
```

Omnibus:	23.070	Durbin-Watson:	1.565
Prob(Omnibus):	0.000	Jarque-Bera (JB):	9.935
Skew:	0.125	Prob(JB):	0.00696
Kurtosis:	2.269	Cond. No.	62.5

```
=====
```



```
lm_0 = smf.ols(formula = "api00 ~ some_col", data = elemapi2.query('yr_rnd == 0')).fit()  
lm_0.summary()
```

#### OLS Regression Results


Dep. Variable:	api00	R-squared:	0.016			
Model:	OLS	Adj. R-squared:	0.013			
Method:	Least Squares	F-statistic:	4.915			
Date:	Mon, 03 Sep 2018	Prob (F-statistic):	0.0274			
Time:	11:01:38	Log-Likelihood:	-1938.2			
No. Observations:	308	AIC:	3880.			
Df Residuals:	306	BIC:	3888.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	655.1103	15.237	42.995	0.000	625.128	685.093
some_col	1.4094	0.636	2.217	0.027	0.158	2.660
Omnibus:	63.461	Durbin-Watson:	1.531			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	13.387			
Skew:	-0.003	Prob(JB):	0.00124			
Kurtosis:	1.979	Cond. No.	48.9			

Some\_col only

# Interaction continuous categorical using yr-rnd 0



```
plt.scatter(elemapi2.query('yr_rnd == 0').some_col.values, elemapi2.query('yr_rnd == 0').api00.values, label = "actual")
plt.scatter(elemapi2.query('yr_rnd == 0').some_col.values, lm_0.predict(), c = "r", label = "predict")
plt.legend()
plt.show()
```






# Interaction continuous categorical using yr-end 1



```
lm_1 = smf.ols(formula = "api00 ~ some_col", data = elemapi2.query('yr_rnd == 1')).fit()
plt.scatter(elemapi2.query('yr_rnd == 1').some_col.values, elemapi2.query('yr_rnd == 1').api00.values, label = "actual")
plt.scatter(elemapi2.query('yr_rnd == 1').some_col.values, lm_1.predict(), c = "r", label = "predict")
plt.legend()
plt.show()
```



The slope of the regression looks much steeper for the year round schools than for the non-year round schools.

(slope for the year round schools to be higher (6.55) than non-year round schools (1.4)). We can compare these to see if these are significantly different from each other by including the



```
yrxsome_elemap_i = elemapi2
yrxsome_elemap_i["yrxsome"] = yrxsome_elemap_i.yr_rnd * yrxsome_elemap_i.some_col

lm = smf.ols("api00 ~ some_col + yr_rnd + yrxsome", data = yrxsome_elemap_i).fit()
lm.summary()
```

OLS Regression Results

Dep. Variable:	api00	R-squared:	0.283
Model:	OLS	Adj. R-squared:	0.277
Method:	Least Squares	F-statistic:	52.05
Date:	Mon, 03 Sep 2018	Prob (F-statistic):	2.21e-28
Time:	11:12:21	Log-Likelihood:	-2483.6
No. Observations:	400	AIC:	4975.
Df Residuals:	396	BIC:	4991.
Df Model:	3		
Covariance Type:	nonrobust		


	coef	std err	t	P> t	[0.025	0.975]
Intercept	655.1103	14.035	46.677	0.000	627.518	682.703
some_col	1.4094	0.586	2.407	0.017	0.258	2.561
yr_rnd	-248.0712	29.859	-8.308	0.000	-306.773	-189.369
yrxsome	5.9932	1.577	3.800	0.000	2.893	9.094

Interaction of some\_col by yr\_rnd  
(interaction of a continuous variable  
by a categorical variable)

Computing interactions manually  
(variable that is the interaction of some college (some\_col)  
and year round schools (yr\_rnd) called **yrxsome**)

```
lt.scatter(yrxsome_elemap_i.query('yr_rnd == 0').some_col, lm.predict()[yrxsome_elemap_i.yr_rnd.values == 0], marker = "+")
lt.scatter(yrxsome_elemap_i.query('yr_rnd == 1').some_col, lm.predict()[yrxsome_elemap_i.yr_rnd.values == 1], c = "r", marker = "x")
```

<matplotlib.collections.PathCollection at 0x1fec843e198>




The two lines have quite different slopes, consistent with the fact that the yrxsome interaction was significant.

## create a plot including the data points



```
l.query('yr_rnd == 0').some_col, lm.predict()[yrxsome_elemapl.yr_rnd.values == 0], marker = "+")  
l.query('yr_rnd == 1').some_col, lm.predict()[yrxsome_elemapl.yr_rnd.values == 1], c = "r", marker = "x")  
l.some_col, yrxsome_elemapl.api00, c = "black", marker = "*")
```

<matplotlib.collections.PathCollection at 0x1fec84a72e8>





Make separate variables for the api00 scores for the two types of schools called api0 for the non-year round schools and api1 for the year round schools.




```
.query('yr_rnd == 0').some_col, lm.predict()[yrxsome_elemapr.yr_rnd.values == 0], marker = "+")  
.query('yr_rnd == 1').some_col, lm.predict()[yrxsome_elemapr.yr_rnd.values == 1], c = "r", marker = "x")  
.query('yr_rnd == 0').some_col, yrxsome_elemapr.query('yr_rnd == 0').api00, c = "black", marker = "*")  
.query('yr_rnd == 1').some_col, yrxsome_elemapr.query('yr_rnd == 1').api00, c = "r", marker = "d")
```

&lt;

&gt;

<matplotlib.collections.PathCollection at 0x1fec850f2e8>



split data to yr\_rnd = 0 group and yr\_rnd = 1 group. Then run regression of api00 to some col in each group separately.



```
yrxsome_elemap_i_0 = yrxsome_elemap_i.query('yr_rnd == 0')
yrxsome_elemap_i_1 = yrxsome_elemap_i.query('yr_rnd == 1')
lm_0 = smf.ols("api00 ~ some_col", data = yrxsome_elemap_i_0).fit()
print (lm_0.summary())
print ('\n')
lm_1 = smf.ols("api00 ~ some_col", data = yrxsome_elemap_i_1).fit()
print (lm_1.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          api00    R-squared:                0.016
Model:                  OLS      Adj. R-squared:           0.013
Method:                 Least Squares    F-statistic:         4.915
Date:                  Mon, 03 Sep 2018    Prob (F-statistic):    0.0274
Time:                  11:23:31    Log-Likelihood:       -1938.2
No. Observations:      308    AIC:                  3880.
Df Residuals:          306    BIC:                  3888.
Df Model:               1
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	655.1103	15.237	42.995	0.000	625.128	685.093
some_col	1.4094	0.636	2.217	0.027	0.158	2.660

```
=====
Omnibus:                 63.461    Durbin-Watson:           1.531
Prob(Omnibus):           0.000    Jarque-Bera (JB):        13.387
Skew:                   -0.003    Prob(JB):                0.00124
Kurtosis:                1.979    Cond. No.                 48.9
=====
```



```
lm = smf.ols("api00 ~ some_col + yr_rnd + yrxsome", data = yrxsome_elemap).fit()
lm.summary()
```



## OLS Regression Results

<b>Dep. Variable:</b>	api00	<b>R-squared:</b>	0.283
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.277
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	52.05
<b>Date:</b>	Mon, 03 Sep 2018	<b>Prob (F-statistic):</b>	2.21e-28
<b>Time:</b>	11:25:56	<b>Log-Likelihood:</b>	-2483.6
<b>No. Observations:</b>	400	<b>AIC:</b>	4975.
<b>Df Residuals:</b>	396	<b>BIC:</b>	4991.
<b>Df Model:</b>	3		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	655.1103	14.035	46.677	0.000	627.518	682.703
<b>some_col</b>	1.4094	0.586	2.407	0.017	0.258	2.561
<b>yr_rnd</b>	-248.0712	29.859	-8.308	0.000	-306.773	-189.369
<b>yrxsome</b>	5.9932	1.577	3.800	0.000	2.893	9.094

<b>Omnibus:</b>	23.863	<b>Durbin-Watson:</b>	1.593
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	9.350

The coefficient for some\_col in the combined analysis is the same as the coefficient for some\_col for the non-year round schools?


This is because non-year round schools are the reference group. Then, the coefficient for the yrxsome interaction in the combined analysis is the Bsome\_col for the year round schools (7.4) minus Bsome\_col for the non year round schools (1.41) yielding 5.99

Show the regression for both types of schools with the interaction term



```
plt.scatter(yrxsome_elemap_i.query('yr_rnd == 0').some_col, lm.predict()[yrxsome_elemap_i.yr_rnd.values == 0], marker = "+")
plt.scatter(yrxsome_elemap_i.query('yr_rnd == 1').some_col, lm.predict()[yrxsome_elemap_i.yr_rnd.values == 1], c = "r", marker = "x")
```

<matplotlib.collections.PathCollection at 0x1fec8587198>



```
lm = smf.ols("api00 ~ some_col + yr_rnd + yr_rnd * some_col ", data = yrxsome_elemap).fit()
lm.summary()
```

#### OLS Regression Results

Dep. Variable:	api00	R-squared:	0.283			
Model:	OLS	Adj. R-squared:	0.277			
Method:	Least Squares	F-statistic:	52.05			
Date:	Mon, 03 Sep 2018	Prob (F-statistic):	2.21e-28			
Time:	11:32:19	Log-Likelihood:	-2483.6			
No. Observations:	400	AIC:	4975.			
Df Residuals:	396	BIC:	4991.			
Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	655.1103	14.035	46.677	0.000	627.518	682.703
some_col	1.4094	0.586	2.407	0.017	0.258	2.561
yr_rnd	-248.0712	29.859	-8.308	0.000	-306.773	-189.369
yr_rnd:some_col	5.9932	1.577	3.800	0.000	2.893	9.094
Omnibus:	23.863	Durbin-Watson:	1.593			

the relationship between some\_col and api00 was significantly stronger than for those from non-year round schools. In general, this type of analysis allows you to test whether the strength of the relationship between two continuous

```
lm = smf.ols("api00 ~ some_col + C(mealcat) + some_col * C(mealcat)", data = elemapi2).fit()
lm.summary()
```

### OLS Regression Results

Dep. Variable:	api00	R-squared:	0.769
Model:	OLS	Adj. R-squared:	0.767
Method:	Least Squares	F-statistic:	263.0
Date:	Mon, 03 Sep 2018	Prob (F-statistic):	4.13e-123
Time:	11:39:42	Log-Likelihood:	-2256.6
No. Observations:	400	AIC:	4525.
Df Residuals:	394	BIC:	4549.
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	825.8937	11.992	68.871	0.000	802.318	849.470
C(mealcat)[T.2]	-239.0300	18.665	-12.806	0.000	-275.725	-202.334
C(mealcat)[T.3]	-344.9476	17.057	-20.223	0.000	-378.483	-311.413
some_col	-0.9473	0.487	-1.944	0.053	-1.906	0.011
some_col:C(mealcat)[T.2]	3.1409	0.729	4.307	0.000	1.707	4.575
some_col:C(mealcat)[T.3]	2.6073	0.896	2.910	0.004	0.846	4.369

The prior examples showed how to do regressions with a continuous variable and a categorical variable that has two levels. How about using a categorical variable with three levels, mealcat

The relationship between some\_col and api00 varied, depending on the level of mealcat. In comparing group 1 with group 2, the coefficient for some\_col was significantly different, but there was no difference in the coefficient for some\_col in comparing groups 2 and 3.

